



Are you ready for the future?



High Performance Solutions, Backend Development and Integration Services, Embedded Software Development, Big Data (Visualization, Architecture, Science), Business Intelligence & Analytics



Objectives

- **Mapreduce in general**
 - Motivation
 - Overview
 - Key, Value paradigm
 - WordCount Example
 - Combiners
 - Example: calculate a mean
 - A Mapreduce Job
- **Mapreduce with Python***
 - WordCount Implementation
 - Debugging and local execution
 - Execution on the cluster.
- **References & Exercises**

Objectives

- **Mapreduce in general**
 - Motivation
 - Overview
 - Key,Value paradigm
 - WordCount Example
 - Combiners
 - Example: calculate a mean
 - A Mapreduce Job
- **Mapreduce with Python***
 - WordCount Implementation
 - Debugging and local execution
 - Execution on the cluster.
- **References & Exercises**



MapReduce - Motivation (1)



- > Processors steadily increased speed but **Disk IO speed “almost remained constant”**.
- > Moore’s Law reaching limits.
- > Nowadays, **disk storage is inexpensive** and many-terabyte problems are common.
- > Supercomputers (**vertical scalability**) are expensive but same amount of processors, RAM and disk over many machines (**horizontal scalability**) is cheap.



Objectives

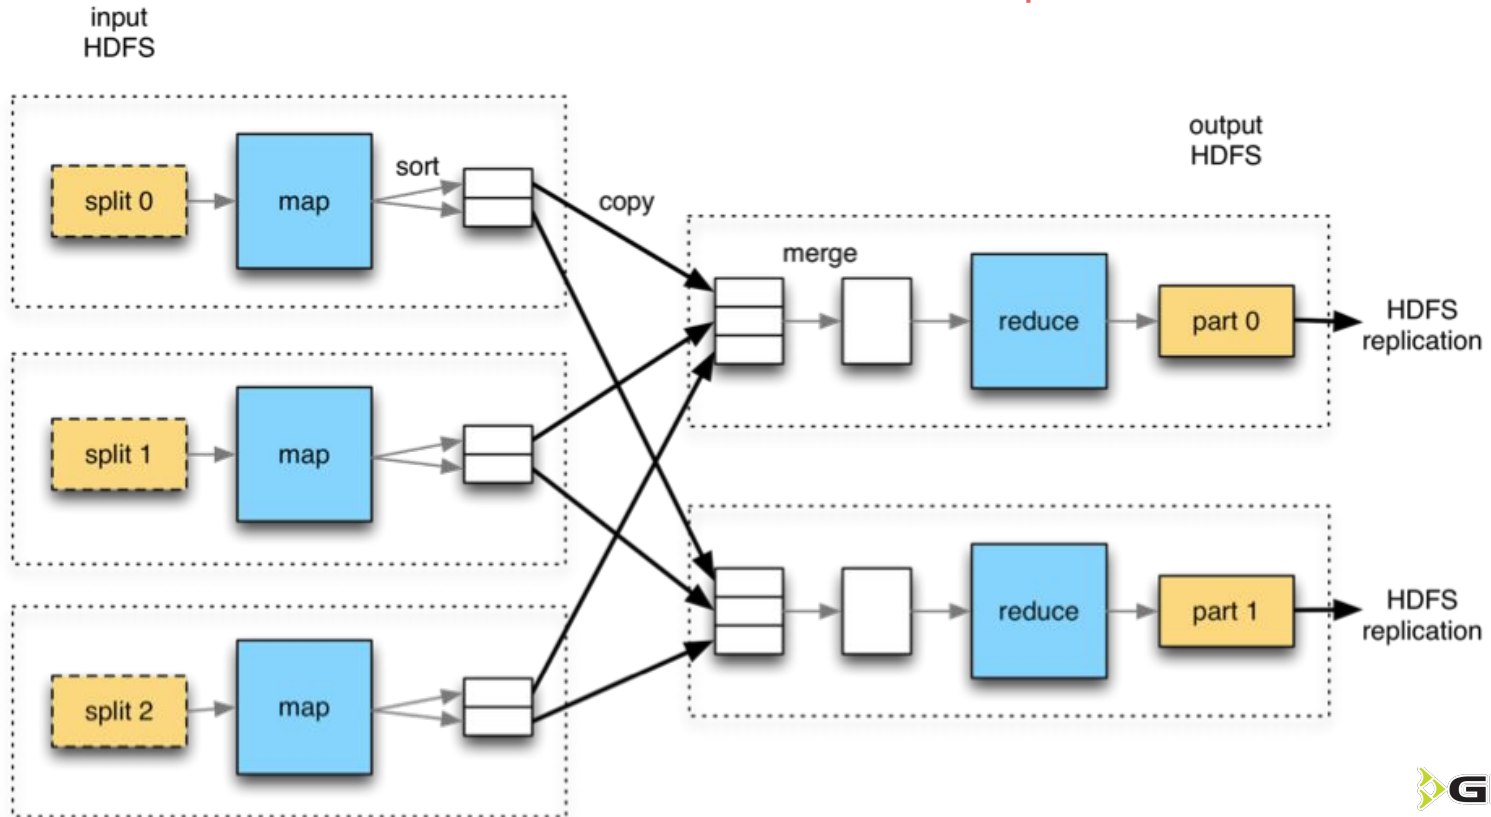
- **Mapreduce in general**
 - Motivation
 - Overview
 - Key, Value paradigm
 - WordCount Example
 - Combiners
 - Example: calculate a mean
 - A Mapreduce Job
- **Mapreduce with Python***
 - WordCount Implementation
 - Debugging and local execution
 - Execution on the cluster.
- **References & Exercises**



MapReduce - Overview (1)

Divide


Conquer





MapReduce - Overview (2)

MapReduce Overview



Objectives

- **Mapreduce in general**
 - Motivation
 - Overview
 - Key, Value paradigm
 - WordCount Example
 - Combiners
 - Example: calculate a mean
 - A Mapreduce Job
- **Mapreduce with Python***
 - WordCount Implementation
 - Debugging and local execution
 - Execution on the cluster.
- **References & Exercises**

MapReduce - Key,Value Paradigm (1)



- > Map and Reduce tasks consume and emit “(key,value)” pairs
- > **keys** and **values** can be of any data type: string, int, **lists**, **custom objects**,...
- > Example:

input: $(k_{\alpha}, \text{value}_{\alpha})$

output: $(k_{\beta}, \text{value}_{\beta})$

Change of alphabets in subindices indicate possible different data types

e.g.,

input: $(5, \text{"this is my pc"})$

output: $(\text{"pc"}, 1)$

input: $(10, \text{"cat"})$

output: $(\text{"letters"}, [\text{"c"}, \text{"a"}, \text{"t"}])$

MapReduce - Key, Value Paradigm (2)



input

(k_{α}, v_{α})

(k_{β}, v_{β})

(k_{γ}, v_{γ})

...

(k_{ω}, v_{ω})

MapReduce - Key, Value Paradigm (3)



input

(k_{α}, v_{α})

(k_{β}, v_{β})


(k_{γ}, v_{γ})

...

(k_{ω}, v_{ω})


Let's forget about these for
a moment

MapReduce - Key, Value Paradigm (4)



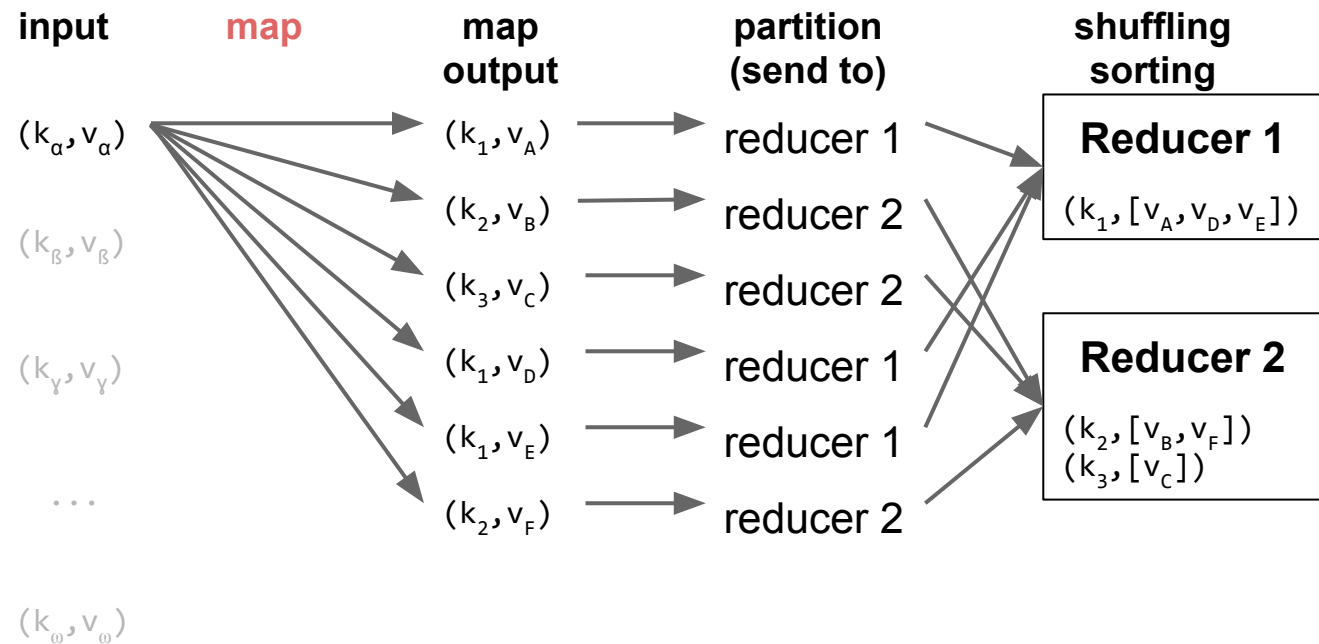
Change of alphabets in subindices indicate possible different data types

MapReduce - Key, Value Paradigm (5)



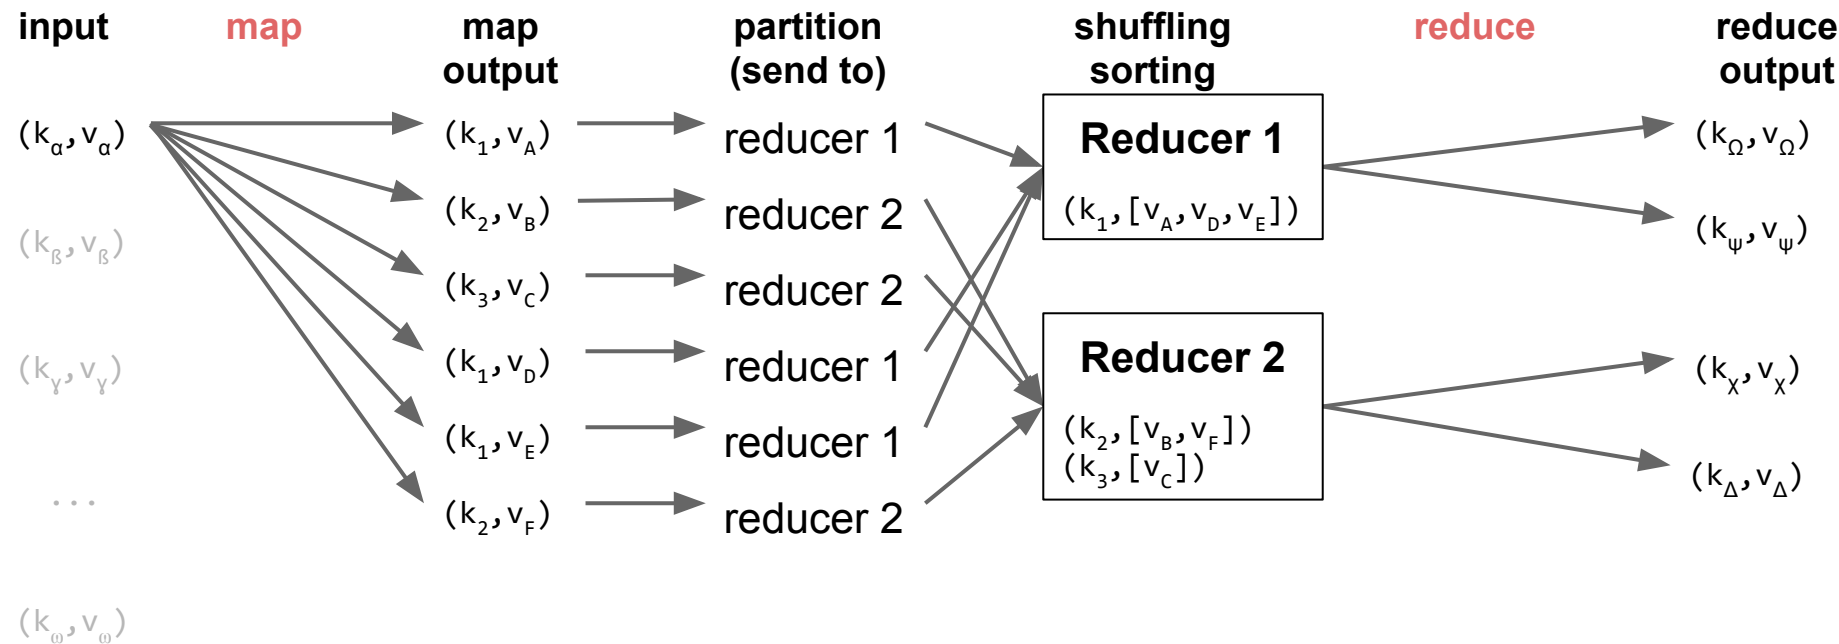
Change of alphabets in subindices indicate possible different data types

MapReduce - Key, Value Paradigm (6)



Change of alphabets in subindices indicate possible different data types

MapReduce - Key,Value Paradigm (7)

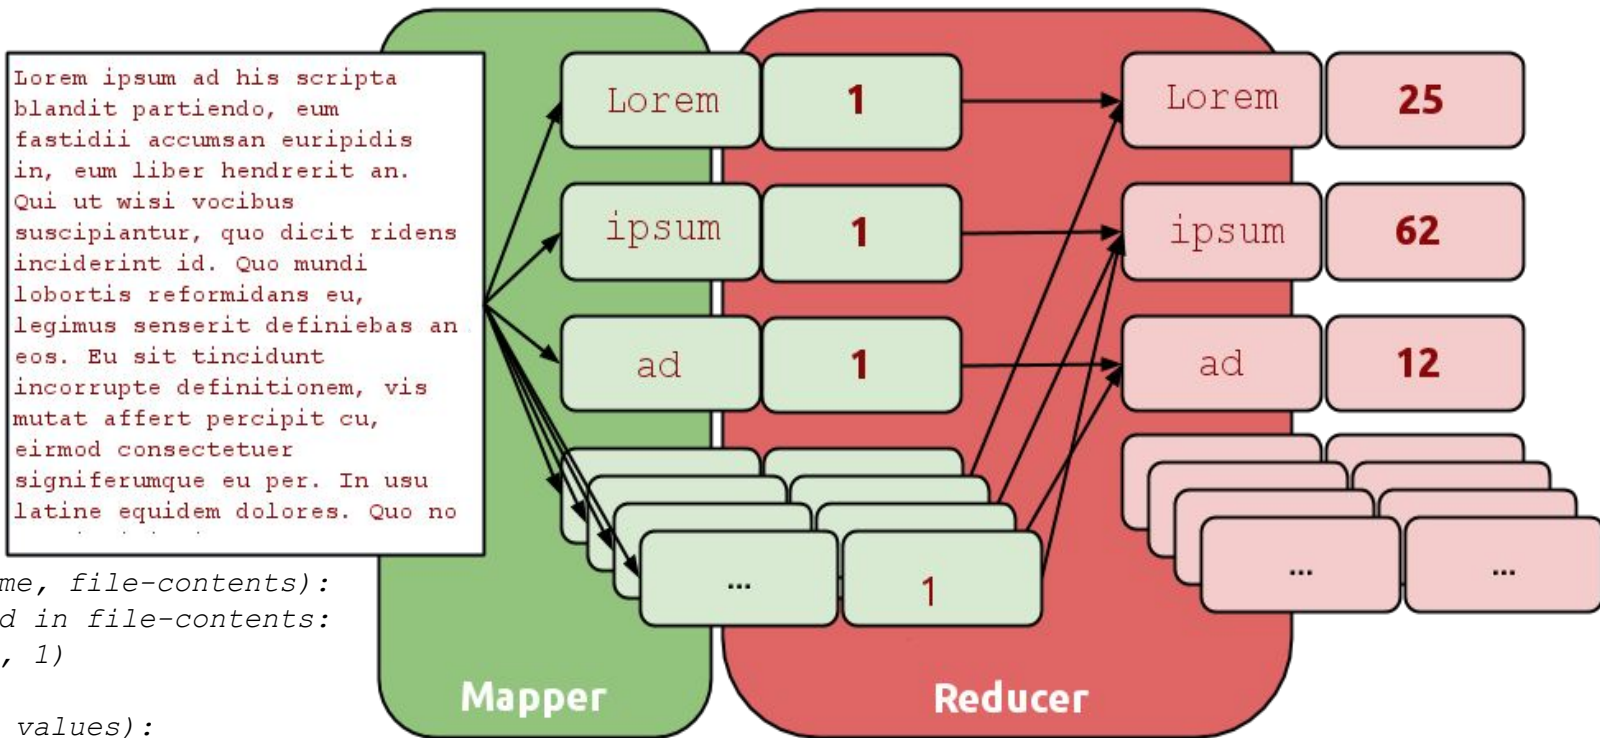


Change of alphabets in subindices indicate possible different data types

Objectives

- **Mapreduce in general**
 - Motivation
 - Overview
 - Key, Value paradigm
 - **WordCount Example**
 - Combiners
 - Example: calculate a mean
 - A Mapreduce Job
- **Mapreduce with Python***
 - WordCount Implementation
 - Debugging and local execution
 - Execution on the cluster.
- **References & Exercises**

MapReduce - WordCount is hadoop's "Hello World"



```
mapper (filename, file-contents):  
    for each word in file-contents:  
        emit (word, 1)
```

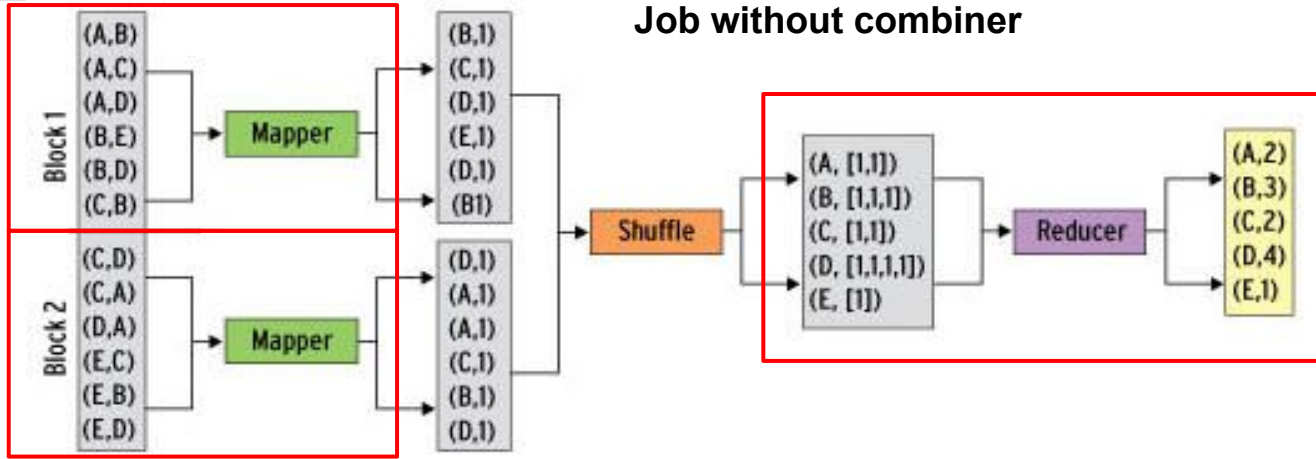
```
reducer (word, values):  
    sum = 0  
    for each value in values:  
        sum = sum + value  
    emit (word, sum)
```

Objectives

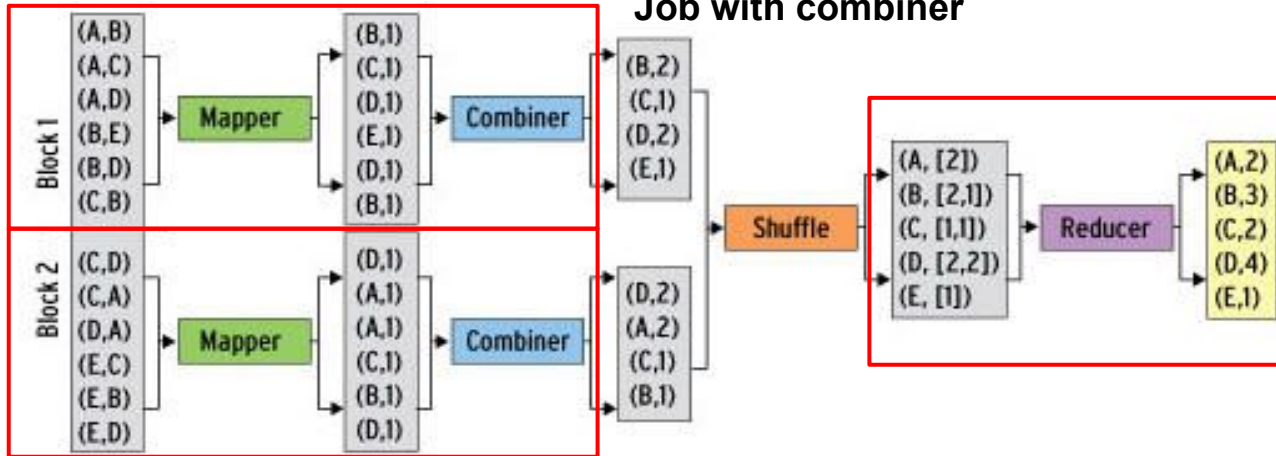
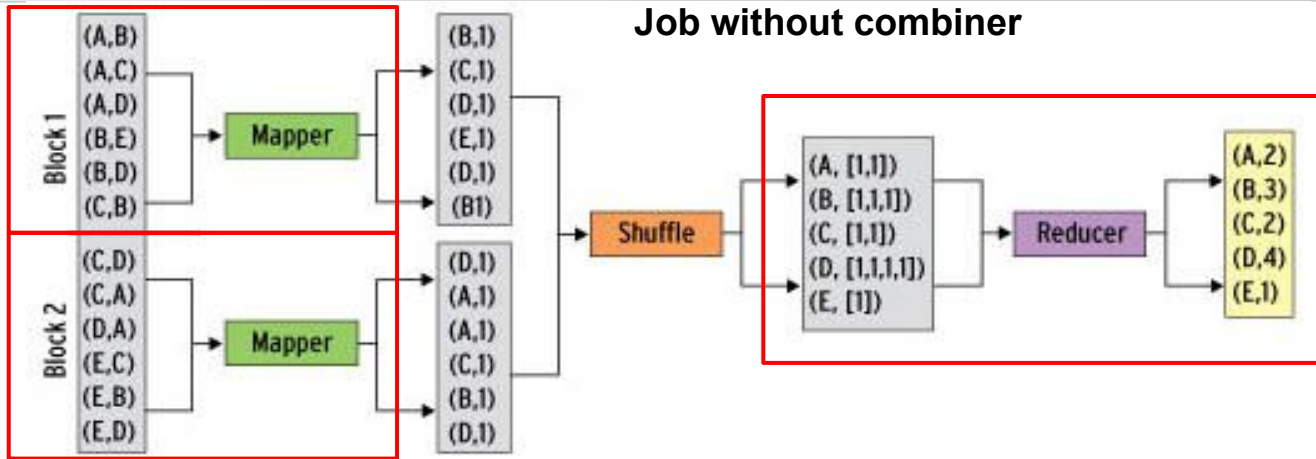
- **Mapreduce in general**
 - Motivation
 - Overview
 - Key,Value paradigm
 - WordCount Example
 - **Combiners**
 - Example: calculate a mean
 - A Mapreduce Job
- **Mapreduce with Python***
 - WordCount Implementation
 - Debugging and local execution
 - Execution on the cluster.
- **References & Exercises**



MapReduce - Jobs with and without Combiner (1)



MapReduce - Jobs with and without Combiner (2)



Data sent
through the network
is much less!
Job significantly faster!

MapReduce - Jobs with and without Combiner (3)



Example,

Say we want to calculate the mean measurement in AR:

node1 mappers outputs:

AR, 500

AR, 1000

AR, 750

node2 mappers outputs:

AR, 500

AR, 800

node3 mappers outputs:

AR, 100

AR, 200

Without Combiner:

$$(500 + 1000 + 750 + 500 + 800 + 100 + 200) / 7 = 550$$

With Combiners (**Common Pitfall**):

$$\begin{aligned} & (500+1000+750)/3 + (500+800)/2 + (100+200)/2 = \\ & (750 + 650 + 150) / 3 = 516,66 \end{aligned}$$

Objectives

- **Mapreduce in general**
 - Motivation
 - Overview
 - Key,Value paradigm
 - WordCount Example
 - Combiners
 - Example: calculate a mean
 - A Mapreduce Job
- **Mapreduce with Python***
 - WordCount Implementation
 - Debugging and local execution
 - Execution on the cluster.
- **References & Exercises**

Task: using mapreduce calculate the mean of the variables height, weight and age (**column-wise**).

txt input file		
Height (cm)	Weight (Kg)	Age (yrs)
180	80	20
160	65	25
175	85	50
192	100	60
181	78	35
150	50	13

$$\bar{x} = \frac{1}{N} \sum_{i=1}^N x_i$$

Idea:

accumulate partial sums and counts, leaving the calculation of the average to the reducer

use as value: <partial sum>_<partial count>

MapReduce - Example: Calculate a mean (2)



Using 3 mappers, **no combiners** and 2 reducers

input splits

```
{0: "180 80 20  
160 65 25  
175 85 50"}
```

```
{28: "192 100 60  
181 78 35"}
```


```
{48: "150 50 13"}
```

$$\bar{x} = \frac{1}{N} \sum_{i=1}^N x_i$$

MapReduce - Example: Calculate a mean (3)



Using 3 mappers, **no combiners** and 2 reducers

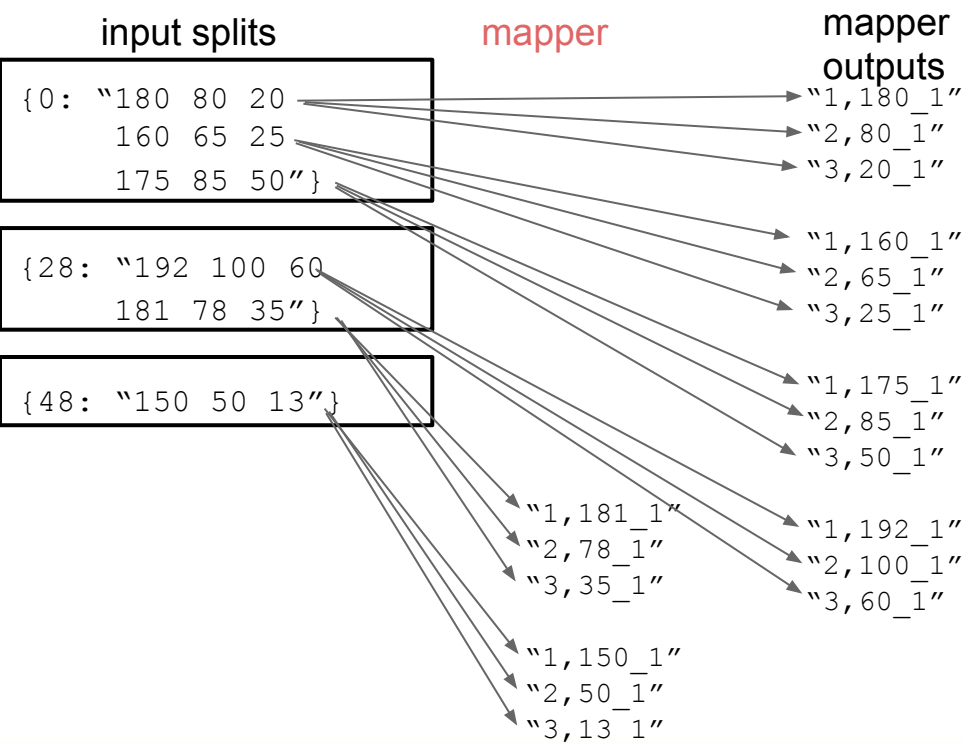


$$\bar{x} = \frac{1}{N} \sum_{i=1}^N x_i$$

MapReduce - Example: Calculate a mean (4)



Using 3 mappers, **no combiners** and 2 reducers



AVOID THIS!!!

mapper's outputs are already larger than input!

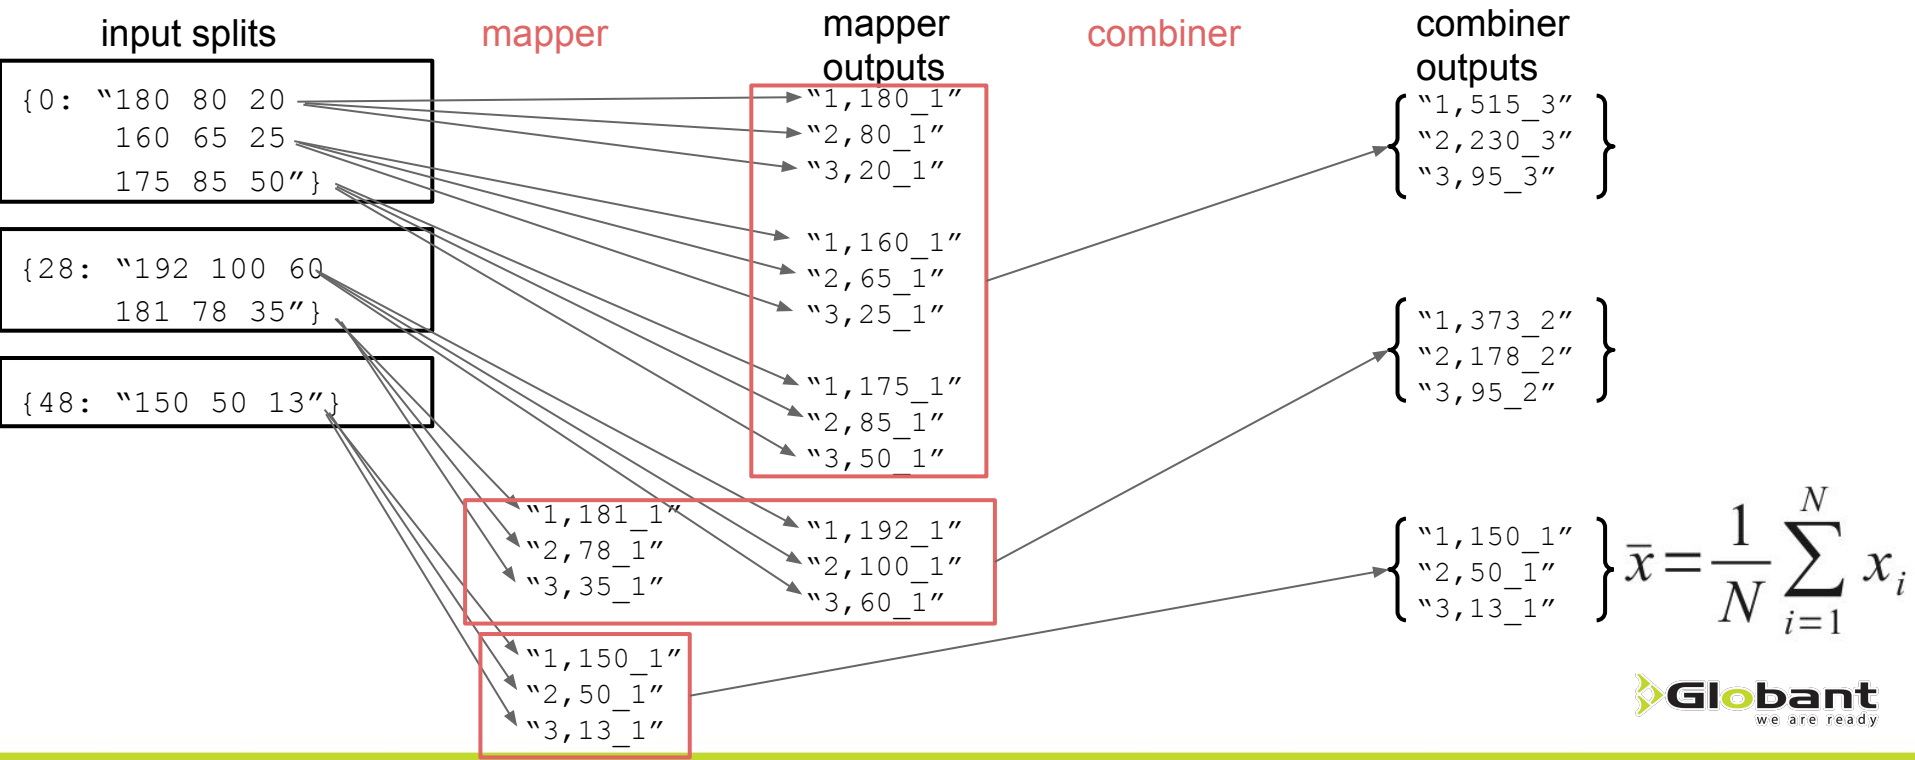
Use combiners to summarize (or accumulators in the mappers if streaming API).

$$\bar{x} = \frac{1}{N} \sum_{i=1}^N x_i$$

MapReduce - Example: Calculate a mean (5)



Using 3 mappers, **combiners** and 2 reducers



MapReduce - Example: Calculate a mean (6)



Using 3 mappers, combiners and 2 reducers

input splits

mapper+combiner

mapper
outputs

{0: "180 80 20
160 65 25
175 85 50"}

{ "1,515_3"
"2,230_3"
"3,95_3" }

{28: "192 100 60
181 78 35"}

{ "1,373_2"
"2,178_2"
"3,95_2" }

{48: "150 50 13"}


{ "1,150_1"
"2,50_1"
"3,13_1" }

$$\bar{x} = \frac{1}{N} \sum_{i=1}^N x_i$$

MapReduce - Example: Calculate a mean (7)



Using 3 mappers, combiners and 2 reducers

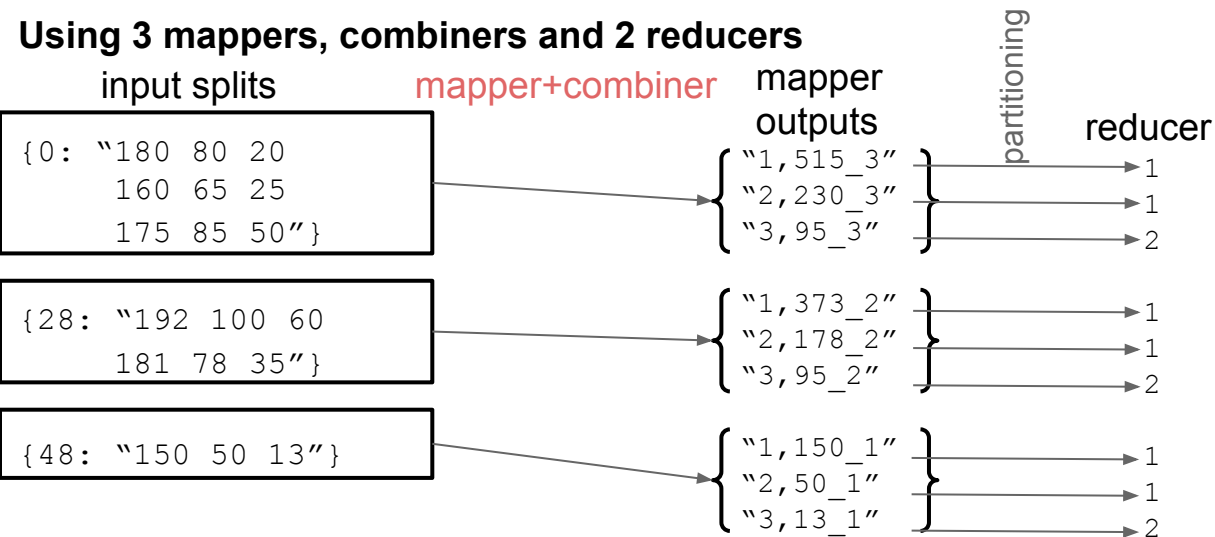


$$\bar{x} = \frac{1}{N} \sum_{i=1}^N x_i$$

MapReduce - Example: Calculate a mean (8)



Using 3 mappers, combiners and 2 reducers



Shuffling and sorting

reducer 1:

{1, [150_1, 373_2, 515_3]}
{2, [178_2, 230_3, 50_1]}

reducer 2:

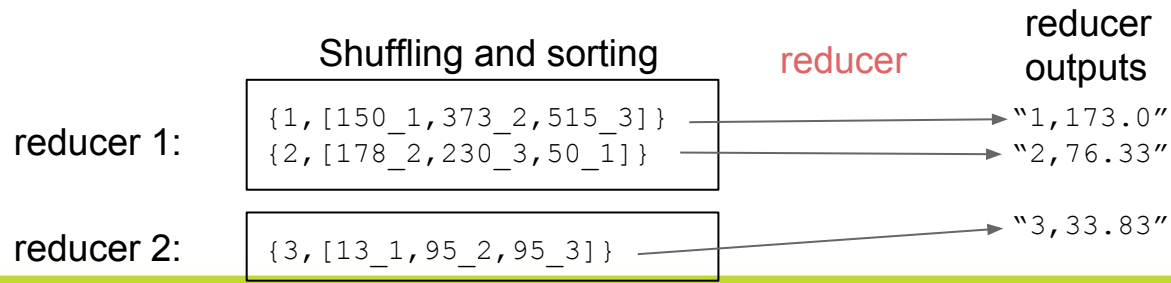

{3, [13_1, 95_2, 95_3]}

$$\bar{x} = \frac{1}{N} \sum_{i=1}^N x_i$$

MapReduce - Example: Calculate a mean (9)



Using 3 mappers, combiners and 2 reducers

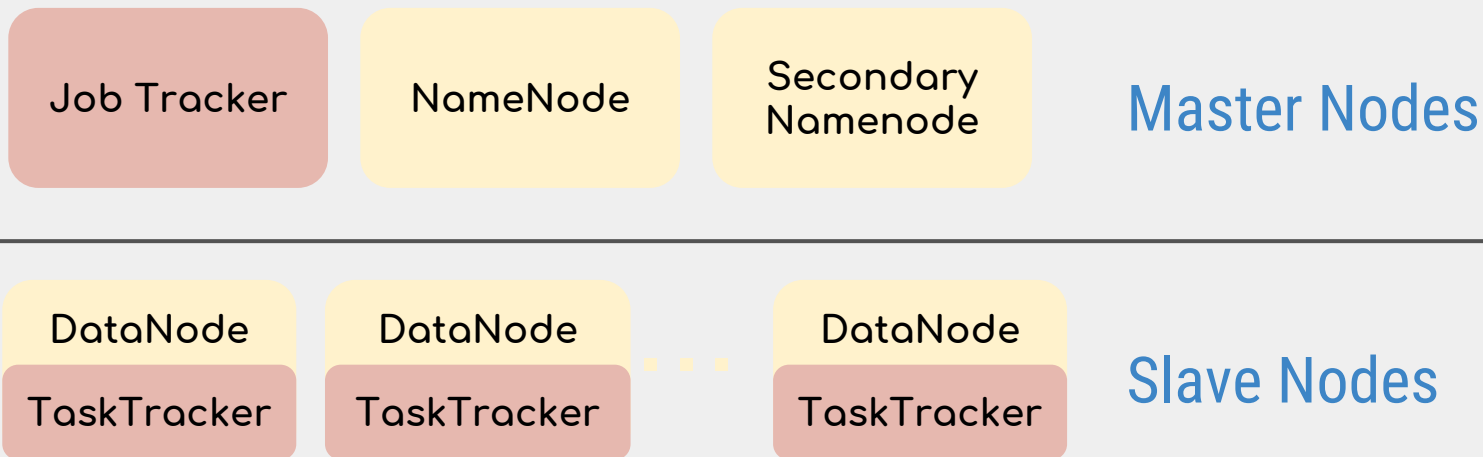


$$\bar{x} = \frac{1}{N} \sum_{i=1}^N x_i$$

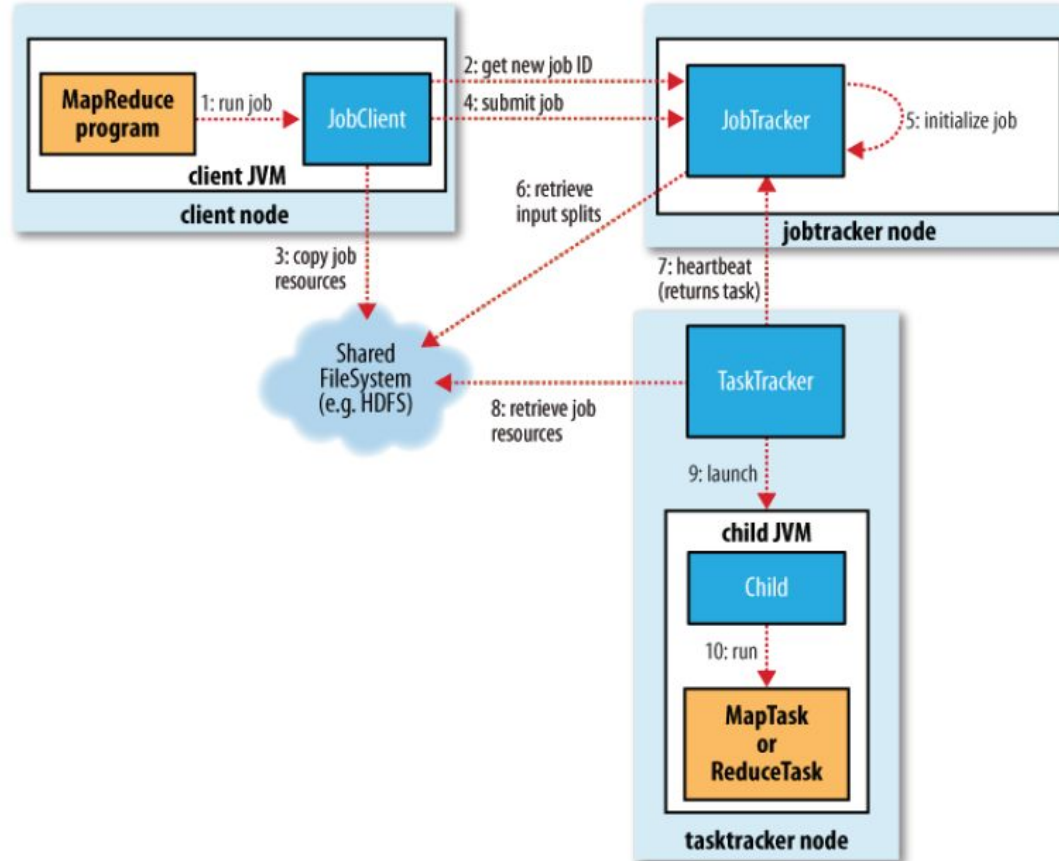
Objectives

- **Mapreduce in general**
 - Motivation
 - Overview
 - Key,Value paradigm
 - WordCount Example
 - Combiners
 - Example: calculate a mean
 - **A Mapreduce Job**
- **Mapreduce with Python***
 - WordCount Implementation
 - Debugging and local execution
 - Execution on the cluster.
- **References & Exercises**

MapReduce - Architecture of a Cluster



MapReduce - How is the code executed?





MapReduce - APIs



- > Streaming API (today)
 - works practically any programming language (**python**, bash, perl, c++, etc.)
 - ideal for prototyping !
 - **mapper script called one time per split.**
 - **reducer input are (key,value) pair sorted by key not (key,list-of-values)**
 - **reducer script called one time for all full pairs of (key, value)**

- > Java API (next session)
 - requires to know Java, hadoop packages, etc.
 - a bit better performance and more control
 - a bit lengthier to debug.
 - **mapper method invoked one time per record.**
 - **reducer method invoked one time per each pair of (key, list of values)**

Objectives

- **Mapreduce in general**
 - Motivation
 - Overview
 - Key, Value paradigm
 - WordCount Example
 - Combiners
 - Example: calculate a mean
 - A Mapreduce Job
- **Mapreduce with Python***
 - WordCount Implementation
 - Debugging and local execution
 - Execution on the cluster.
- **References & Exercises**



MapReduce - Streaming API (python, php, c++, etc)



- > This API allows practically **any programming language** (we'll focus on python)
- > (Key,Value) pairs fed as a **continuous stream of text** with sorted keys.
- > Data flows into mapper/reducer via *NIX pipes behind the scenes.
- > the Developer must:
 - **read** data from **STDIN** (standard input)
 - **define** key,value **separator**
 - **split and parse** key and value strings **to the correct data type**
 - **identify** when keys change in the stream
 - **emit** by writing to **STDOUT** (standard output)

Objectives

- **Mapreduce in general**
 - Motivation
 - Overview
 - Key, Value paradigm
 - WordCount Example
 - Combiners
 - Example: calculate a mean
 - A Mapreduce Job
- **Mapreduce with Python***
 - WordCount Implementation
 - Debugging and local execution
 - Execution on the cluster.
- **References & Exercises**

```
mapper (filename, file-contents):  
    for each word in file-contents:  
        emit (word, 1)
```

```
reducer (word, values):  
    sum = 0  
    for each value in values:  
        sum = sum + value  
    emit (word, sum)
```


MapReduce - wordcount_mapper.py



```
#!/usr/bin/env python
import sys

# input comes from STDIN (standard input)
for line in sys.stdin:
    # remove leading and trailing whitespace
    line = line.strip()
    # split the line into a list of words
    words = line.split(' ')
    # increase counters
    for word in words:
        # write the results to STDOUT (standard output);
        # what we output here will be the input for the
        # Reduce step, tab-delimited.
        # The trivial word count is 1
        print '%s\t%s' % (word, 1)
```

MapReduce - wordcount_reducer.py



```
#!/usr/bin/env python
import sys
current_word = None
current_count = 0
word = None
# input comes from STDIN
for line in sys.stdin:
    # remove leading and trailing whitespace
    line = line.strip()
    # parse the input we got from wordcount_mapper.py
    word, count = line.split('\t', 1)
    # convert count (currently a string) to int
    count = int(count)
    if current_word == word:
        current_count += count
    else:
        if current_word:
            # write result to STDOUT
            print '%s\t%s' % (current_word, current_count)
            current_count = count
            current_word = word

# do not forget to output the last word if needed!
if current_word == word:
    print '%s\t%s' % (current_word, current_count)
```

Objectives

- **Mapreduce in general**
 - Motivation
 - Overview
 - Key,Value paradigm
 - WordCount Example
 - Combiners
 - Example: calculate a mean
 - A Mapreduce Job
- **Mapreduce with Python***
 - WordCount Implementation
 - Debugging and local execution
 - Execution on the cluster.
- **References & Exercises**

Testing Execution without Hadoop!! (useful to debug)

Run from shell:

```
cat <input txt file(s)> | python ./wordcount_mapper.py | sort -k1,1 | python ./wordcount_reducer.py
```

The output is shown directly on the screen:

```
"          28
"'A          1
"'About      1
"'Absolute    1
"'Ah!'        2
"'Ah,         2
"'Ample.'      1
"'And        10
"'Are         1
"'Arthur!'     1
"'As          1
...

```



This **only works** if the code is totally independent of the **hadoop environment**, e.g., no use of environment variables, etc.

Execution with Hadoop using the VM (pseudo cluster)

After copying the code and the sample data to the VM, execute from shell:

```
hadoop \  
jar <path-to-hadoop-parent-dir>/hadoop-2.6.0/share/hadoop/tools/lib/hadoop-streaming-2.6.0.jar \  
-D mapred.map.tasks=5 \      ←optional  
-D mapred.reduce.tasks=2 \   ←optional  
-files <code-dir>/wordcount_mapper.py,<code-dir>/wordcount_reducer.py \  
-mapper "python ./wordcount_mapper.py" \  
-reducer "python ./wordcount_reducer.py" \  
-input <hdfs-input-dir> \  
-output <hdfs-output-dir>
```

The output is written to the HDFS files:

```
<hdfs-output-dir>/part-00000  
<hdfs-output-dir>/part-00001
```

To see a part of one of the files:

```
hadoop fs -tail <hdfs-output-dir>/part-00000
```

Objectives

- **Mapreduce in general**
 - Motivation
 - Overview
 - Key,Value paradigm
 - WordCount Example
 - Combiners
 - Example: calculate a mean
 - A Mapreduce Job
- **Mapreduce with Python***
 - WordCount Implementation
 - Debugging and local execution
 - Execution on the cluster.
- **References & Exercises**

Execution on the AWS cluster

After copying the python code to the main server, execute from shell:

```
hadoop \  
jar /home/hadoop/hadoop/hadoop-2.5.2/share/hadoop/tools/lib/hadoop-streaming-2.5.2.jar \  
-D mapred.map.tasks=5 \      ←optional  
-D mapred.reduce.tasks=2 \   ←optional  
-files <code-dir>/wordcount_mapper.py,<code-dir>/wordcount_reducer.py \  
-mapper "python ./wordcount_mapper.py" \  
-reducer "python ./wordcount_reducer.py" \  
-input /user/hadoop/mapreduce/data/books \  
-output <hdfs-output-dir>
```

The output is written to the HDFS files:

```
<hdfs-output-dir>/part-00000  
<hdfs-output-dir>/part-00001
```

To see a part of one of the files:

```
hadoop fs -tail <hdfs-output-dir>/part-00000
```

Objectives

- **Mapreduce in general**
 - Motivation
 - Overview
 - Key, Value paradigm
 - WordCount Example
 - Combiners
 - Example: calculate a mean
 - A Mapreduce Job
- **Mapreduce with Python***
 - WordCount Implementation
 - Debugging and local execution
 - Execution on the cluster.
- **References & Exercises**



Mapreduce - References



- > Yahoo Mapreduce tutorial

<https://developer.yahoo.com/hadoop/tutorial/module4.html>

- > White, Tom. Hadoop: The definitive guide (a.k.a. "The Elephant Book")

- > Apache documentation for Streaming API

<http://hadoop.apache.org/docs/stable/hadoop-mapreduce-client/hadoop-mapreduce-client-core/HadoopStreaming.html>



Mapreduce - Exercises



- > Avoid peeking at the solutions on the web !
- > Use WordCount code as starting point.
- > Don't code and debug right off using the full input dataset. Subset or **create a simpler dataset first (better)**.
- > Focus on the mapreduces first. If any other "plumbing code" is needed, do it after that and use a manual solution in the meantime.
- > **Don't share coded solutions, instead share hints on how to proceed and give your peers the opportunity to learn for themselves.**



Mapreduce - Exercise 1



Inverted Index

Given a directory with books in txt format, write a mapreduce which outputs an inverted index, i.e., a table that associates a word with the books and the corresponding positions at which it occurs (http://en.wikipedia.org/wiki/Inverted_index).

Dataset URL: [here](#)

HDFS Cluster DataSet path: `/user/hadoop/mapreduce/data/books`

hint 1: Suggested output example (not real data):

Love `alice_in_wonderland.txt:100,the_prince.txt:900,the_prince.txt:1050`

hint 2: Given the mapper doesn't receive the filename as input. A Hadoop Configured Parameter (environment variable) could help to retrieve the filename from which the word comes.



Mapreduce - Exercise 2



Column-wise Variance of a matrix

Given a csv file without headers, calculate the sample variance (s^2) of each column.

(<http://en.wikipedia.org/wiki/Variance>)

HDFS DataSet path: `/user/hadoop/mapreduce/data/matrix`

!!!DON'T DOWNLOAD IT FROM THE CLUSTER \$\$\$\$!!!

hint 1: Suggested output: `columnIndex<tab>sampleVariance`.

Example:

```
0    135.6
1     2.2
2   536.9
```

...

hint 2: Assume the file has only numeric values and no entries are missing (no NULLs or empty).

hint 3: The python modules `os` and `sys` will come quite handy.

hint 4: Focus on the the mapreduce. (If any) other plumbing code is required, do it manually, afterwards implement it if you have the time (with python of course!).

$$s^2 = \frac{1}{(N-1)} \sum_{i=1}^N (x_i - \bar{x})^2$$



Mapreduce - Exercise 2



Column-wise Variance of a matrix

question 1:

Assuming your matrix is very big (both rows and columns), how many reducers does it make sense to have?

question 2:

How many mapreduce jobs do you need?

question 3:

Implement combiners and notice the savings in bytes transferred to the reducers.

