



Are you ready for the future?



High Performance Solutions, Backend Development and Integration Services, Embedded
Software Development, Big Data (Visualization, Architecture, Science), Business
Intelligence & Analytics



Objectives

- **Mapreduce APIs**
- **Basics**
 - **Mapreduce Java API**
 - **Maven Dependencies**
 - **WordCount - New Java API**
 - **WordCount - Old Java API**
 - **Execution on the cluster**
 - **MRUnit**
 - **Hadoop IO data Types (Writable)s**
 - **Writing a Writable custom Class**
- **Advanced**
 - **How does the data flow through the tasks**
- **References and Exercises**

Objectives

- **Mapreduce APIs**
- **Basics**
 - Mapreduce Java API
 - Maven Dependencies
 - WordCount - New Java API
 - WordCount - Old Java API
 - Execution on the cluster
 - MRUnit
 - Hadoop IO data Types (Writable)s
 - Writing a Writable custom Class
- **Advanced**
 - How does the data flow through the tasks
- **References and Exercises**



MapReduce APIs



- > Streaming API (last session)
 - works practically any programming language (**python**, bash, perl, c++, etc.)
 - ideal for prototyping !
 - **mapper script called one time per split.**
 - **reducer inputs are (key,value) pairs sorted by key, not (key,list-of-values)**
 - **reducer script called one time for all full pairs of (key, value)**

- > Java API (today)
 - requires to know Java, hadoop packages, etc.
 - a bit better performance and more control
 - a bit lengthier to debug.
 - **mapper method invoked one time per record.**
 - **reducer method invoked one time per each pair of (key, list-of-values)**

Objectives

- **Mapreduce APIs**
- **Basics**
 - **Mapreduce Java API**
 - **Maven Dependencies**
 - **WordCount - New Java API**
 - **WordCount - Old Java API**
 - **Execution on the cluster**
 - **MRUnit**
 - **Hadoop IO data Types (Writable)s**
 - **Writing a Writable custom Class**
- **Advanced**
 - **How does the data flow through the tasks**
- **References and Exercises**



MapReduce Java API



- > Two Mapreduce Java API versions: **Old API** (up to 0.20) and **New API** (after 0.21)
- > Two Mapreduce Framework implementations: **Mapreduce 1** (Classic) and **Mapreduce 2** (YARN)
- > **Both APIs work on both Framework implementations** (We'll focus on NEW API).
- > **(Key,Value) pairs fed one at a time** automatically by framework to mapper/reducer methods.
- > The Developer must (minimally):
 - write **Mapper class** with **map()** method
 - write **Reducer class** with **reduce()** method
 - write **Java Driver class** with **run()** method and **static main method**
- > The framework **serializes/deserializes** built-in data types (using the Writable interface).

Objectives

- **Mapreduce APIs**
- **Basics**
 - **Mapreduce Java API**
 - **Maven Dependencies**
 - **WordCount - New Java API**
 - **WordCount - Old Java API**
 - **Execution on the cluster**
 - **MRUnit**
 - **Hadoop IO data Types (Writable)s**
 - **Writing a Writable custom Class**
- **Advanced**
 - **How does the data flow through the tasks**
- **References and Exercises**

Maven Dependencies (add to pom.xml)



...
<dependencies>

<dependency>

<groupId>org.apache.hadoop</groupId>

<artifactId>hadoop-client</artifactId>

<version>2.5.2</version> ← (the Cluster has Apache Hadoop 2.5.2 and the VM has 2.6.0)

<scope>provided</scope>

</dependency>

<dependency>

<groupId>org.apache.mrunit</groupId>

<artifactId>mrunit</artifactId>

<version>1.1.0</version>

<scope>test</scope>

<classifier>hadoop2</classifier>

</dependency>

</dependencies>

...

Objectives

- **Mapreduce APIs**
- **Basics**
 - Mapreduce Java API
 - Maven Dependencies
 - **WordCount - New Java API**
 - WordCount - Old Java API
 - Execution on the cluster
 - MRUnit
 - Hadoop IO data Types (Writable)s
 - Writing a Writable custom Class
- **Advanced**
 - How does the data flow through the tasks
- **References and Exercises**



WordCount pseudocode



```
mapper (filename, file-contents):  
    for each word in file-contents:  
        emit (word, 1)
```

```
reducer (word, values):  
    sum = 0  
    for each value in values:  
        sum = sum + value  
    emit (word, sum)
```

WordCountMapper.java (New API)



```
package com.hadoop.example.wordcount;

import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;
import java.io.IOException;

public class WordCountMapper extends Mapper<LongWritable, Text, Text, LongWritable> {
    private static final LongWritable ONE = new LongWritable( 1 );

    @Override
    public void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {
        String line = value.toString().toLowerCase();
        String[] words = line.split( " " );
        for (String word : words ) {
            context.write( new Text(word), ONE );
        }
    }
}
```

WordCountReducer.java (New API)



```
package com.hadoop.example.wordcount;

import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;
import java.io.IOException;

public class WordCountReducer extends Reducer<Text, LongWritable, Text, LongWritable> {

    private LongWritable result = new LongWritable();

    @Override
    public void reduce(Text key, Iterable<LongWritable> values, Context context) throws IOException, InterruptedException {
        int sum = 0;
        for (LongWritable value : values) {
            sum += value.get();
        }
        result.set(sum);
        context.write(key, result);
    }
}
```

WordCountDriver.java - part 1 (New API)



```
package com.hadoop.example.wordcount;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.conf.Configured;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.util.Tool;
import org.apache.hadoop.util.ToolRunner;
```

WordCountDriver.java - part 2 (New API)



```
public class WordCountDriver extends Configured implements Tool {
    public static void main(String[] args) throws Exception {
        int res = ToolRunner.run( new Configuration(), new WordCountDriver(), args );
        System.exit( res );
    }

    @Override
    public int run(String[] args) throws Exception {
        if (args.length != 2) {
            System.out.println("Usage: hadoop jar hadoop-wordcount-example-1.0-SNAPSHOT-job.jar"
                               + " [generic options] <in> <out>");

            return 1;
        }
        Job job = new Job( getConf(), "WordCount" );
        job.setJarByClass( getClass() );
        job.setMapperClass( WordCountMapper.class );
        job.setCombinerClass( WordCountReducer.class );
        job.setReducerClass( WordCountReducer.class );
        job.setOutputKeyClass( Text.class );
        job.setOutputValueClass( LongWritable.class );
        FileInputFormat.addInputPath( job, new Path(args[0]) );
        FileOutputFormat.setOutputPath( job, new Path(args[1]) );
        boolean success = job.waitForCompletion( true );
        return success ? 0 : 1;
    }
}
```

Objectives

- Mapreduce APIs
- Basics
 - Mapreduce Java API
 - Maven Dependencies
 - WordCount - New Java API
 - **WordCount - Old Java API**
 - Execution on the cluster
 - MRUnit
 - Hadoop IO data Types (Writable)s
 - Writing a Writable custom Class
- Advanced
 - How does the data flow through the tasks
- References and Exercises

WordCountMapper.java (Old API)



```
package com.hadoop.example.wordcount;

import java.io.IOException;

import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;           // MapreduceBase and Mapper merged in new API as class Mapper
import org.apache.hadoop.mapred.Mapper;                  // new API uses package org.apache.hadoop.mapreduce
import org.apache.hadoop.mapred.OutputCollector;         // OutputCollector and Reporter merged in new API as Context
import org.apache.hadoop.mapred.Reporter;

public class WordCountMapper extends MapReduceBase implements Mapper<LongWritable,Text,Text,LongWritable>{

    private static final LongWritable ONE = new LongWritable( 1 );

    public void map(LongWritable key, Text value, OutputCollector<Text,LongWritable> output, Reporter reporter) throws
    IOException{

        String line = value.toString();
        String[] words = line.split( " " );
        for (String word: words) {
            output.collect(new Text(word), ONE );
        }
    }
}
```


WordCountReducer.java (Old API)



```
package com.hadoop.example.wordcount;

import java.io.IOException;
import java.util.Iterator;

import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;           // MapreduceBase and Reducer merged in new API as class Reducer
import org.apache.hadoop.mapred.Reducer;                // new API uses package org.apache.hadoop.mapreduce
import org.apache.hadoop.mapred.OutputCollector;         // OutputCollector and Reporter merged in new API as Context
import org.apache.hadoop.mapred.Reporter;

public class WordCountReducer extends MapReduceBase implements Reducer<Text, LongWritable, Text, LongWritable> {

    public void reduce(Text key, Iterator values, OutputCollector<Text, LongWritable> output, Reporter reporter) throws
    IOException {
        long sum = 0;
        while( values.hasNext() ) {
            sum += values.next().get();
        }
        output.collect(key, new LongWritable(sum));
    }
}
```

WordCountDriver.java - part 1 (Old API)



```
package com.hadoop.example.wordcount;
import java.io.IOException;

import org.apache.hadoop.conf.Configured;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.FileInputFormat;           // new API uses package org.apache.hadoop.mapreduce
import org.apache.hadoop.mapred.FileOutputFormat;
import org.apache.hadoop.mapred.JobClient;                 // JobClient and JobConf merged in new API as class Job
import org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.util.Tool;
import org.apache.hadoop.util.ToolRunner;
```

WordCountDriver.java - part 2 (Old API)



```
public class WordCountDriver extends Configured implements Tool {

    public static void main(String[] args) throws Exception{
        int exitCode=ToolRunner.run(new WordCountDriver(), args);
        System.exit(exitCode);
    }

    public int run(String[] args) throws IOException{
        Path inputPath = new Path( args[0] );
        Path outputPath = new Path( args[1] );
        JobConf conf = new JobConf( WordCount.class );
        conf.setJobName( "WordCount" );
        FileInputFormat.setInputPaths( conf, inputPath );
        FileOutputFormat.setOutputPath( conf, outputPath );
        conf.setMapperClass( WordCountMapper.class );
        conf.setReducerClass( WordCountReducer.class );
        conf.setMapOutputKeyClass( Text.class );
        conf.setMapOutputValueClass( LongWritable.class );
        conf.setOutputKeyClass( Text.class );
        conf.setOutputValueClass( LongWritable.class );
        JobClient.runJob( conf );
        return 0;
    }
}
```

Objectives

- **Mapreduce APIs**
- **Basics**
 - Mapreduce Java API
 - Maven Dependencies
 - WordCount - New Java API
 - WordCount - Old Java API
 - **Execution on the cluster**
 - MRUnit
 - Hadoop IO data Types (Writable)s
 - Writing a Writable custom Class
- **Advanced**
 - How does the data flow through the tasks
- **References and Exercises**

WordCount Execution on cluster



Compile and build the jar file using maven

```
$ mvn clean package
```

copy the jar file to the cluster

```
$ scp -i ./Training.pem \  
target/hadoop-wordcount-example-1.0-SNAPSHOT-job.jar \  
hadoop@54.172.223.97:~/mapreduce/juan.gaviria
```

Connect to the cluster

```
$ ssh -i ./Training.pem hadoop@54.172.223.97
```

Execute the mapreduce

```
$ hadoop \  
jar ~/mapreduce/juan.gaviria/hadoop-wordcount-example-1.0-SNAPSHOT-job.jar \  
com.hadoop.example.wordcount.WordCountDriver \  
-D mapred.map.tasks=5 \                               ←optional  
-D mapred.reduce.tasks=2 \                             ←optional  
/user/hadoop/mapreduce/data/books \                   ←hdfs input dir  
/user/hadoop/mapreduce/juan.gaviria/<hdfs-output-dir> ←hdfs output dir
```

Objectives

- **Mapreduce APIs**
- **Basics**
 - Mapreduce Java API
 - Maven Dependencies
 - WordCount - New Java API
 - WordCount - Old Java API
 - Execution on the cluster
 - MRUnit
 - Hadoop IO data Types (Writable)s
 - Writing a Writable custom Class
- **Advanced**
 - How does the data flow through the tasks
- **References and Exercises**



WordCountMapReduceTest.java - part 1 (New API)



```
package com.hadoop.example.wordcount;

import java.io.IOException;
import java.util.ArrayList;
import java.util.List;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.MapDriver;
import org.apache.hadoop.mapreduce.ReduceDriver;
import org.junit.Before;
import org.junit.Test;

public class WordCountMapReduceTest {
    MapDriver<LongWritable, Text, Text, LongWritable> mapDriver;
    ReduceDriver<Text, LongWritable, Text, LongWritable> reduceDriver;

    private static final LongWritable ONE = new LongWritable( 1 );

    @Before
    public void setUp() {
        WordCountMapper mapper = new WordCountMapper();
        WordCountReducer reducer = new WordCountReducer();
        mapDriver = MapDriver.newMapDriver(mapper);
        reduceDriver = ReduceDriver.newReduceDriver(reducer);
    }
}
```



WordCountMapReduceTest.java - part 2 (New API)



```
@Test
public void testMapper() throws IOException {
    LongWritable inKey = new LongWritable( 0 );
    Text inValue = new Text( "bar foo bar" );
    mapDriver.withInput( inKey, inValue );

    Text outKey1= new Text( "bar" );
    LongWritable outValue1 = ONE;
    Text outKey2= new Text( "foo" );
    LongWritable outValue2 = ONE;
    Text outKey3= new Text( "bar" );
    LongWritable outValue3 = ONE;
    mapDriver.withOutput( outKey1, outValue1 );
    mapDriver.withOutput( outKey2, outValue2 );
    mapDriver.withOutput( outKey3, outValue3 );
    mapDriver.runTest( true );
}
```




WordCountMapReduceTest.java - part 3 (New API)



```
@Test
public void testReducer() throws IOException {
    Text inKey = new Text( "bar" );
    List<LongWritable> inListValues = new ArrayList<LongWritable>();
    inListValues.add( ONE );
    inListValues.add( ONE );
    reduceDriver.withInput( inKey, inListValues );

    Text outKey = new Text( inKey.toString() );
    LongWritable outValue = new LongWritable( 2 );
    reduceDriver.withOutput( outKey, outValue );
    reduceDriver.runTest();
}
}
```

Objectives

- **Mapreduce APIs**
- **Basics**
 - Mapreduce Java API
 - Maven Dependencies
 - WordCount - New Java API
 - WordCount - Old Java API
 - Execution on the cluster
 - MRUnit
 - **Hadoop IO data Types (Writable)s**
 - Writing a Writable custom Class
- **Advanced**
 - How does the data flow through the tasks
- **References and Exercises**



Hadoop IO Data Types (1)



- > “Writables” are built-in datatypes in `org.apache.hadoop.io` optimized for serialization.
- > Any serialization frameworks can be used, e.g., Avro (<http://avro.apache.org/>) is very popular.
- > The framework **serializes/deserializes** types using the **Writable** interface:

```
package org.apache.hadoop.io;

import java.io.DataOutput;
import java.io.DataInput;
import java.io.IOException;

public interface Writable {
    void write(DataOutput out) throws IOException;    ← Writes to binary stream
    void readFields(DataInput in) throws IOException; ← Reads from binary stream
}
```



Hadoop IO Data Types (2)



- > The framework **sorts keys** using the good old **Comparable** interface:

```
package java.lang;

public interface Comparable<T> {
    int compareTo(T in);      ← Reads from binary stream
}
```

- > Also needs to override **equals()**, **hashCode()** and **toString()** from **java.lang.Object**
- > A **default constructor** is needed too.
- > The framework defines the **WritableComparable** subinterface:

```
package org.apache.hadoop.io;

public interface WritableComparable<T> extends Writable, Comparable<T> {
}
```

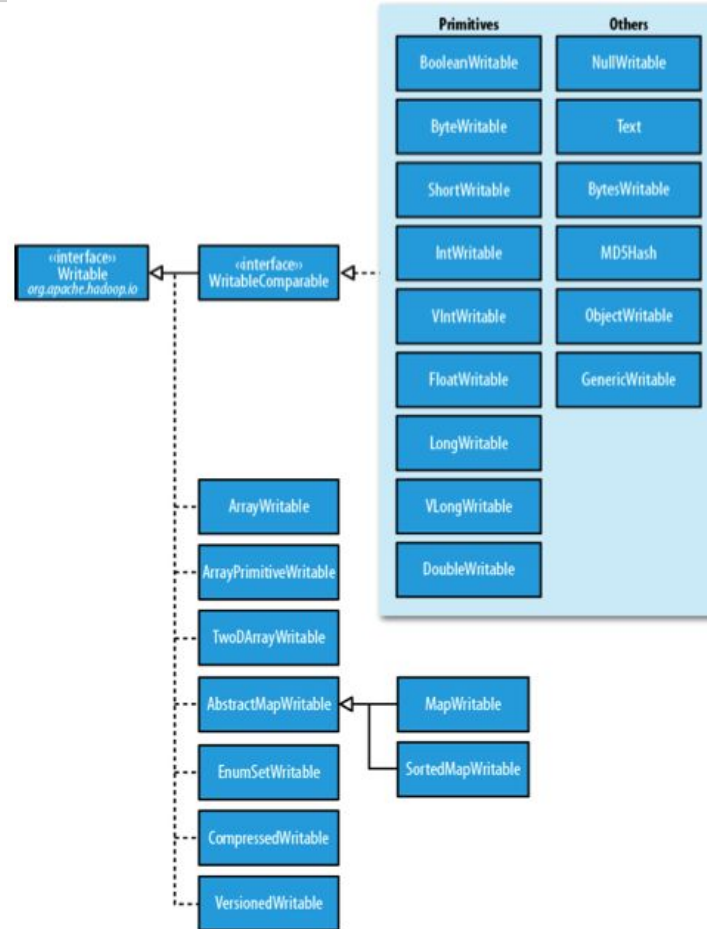


Hadoop IO Data Types (3)

- > Hierarchy of **Writable** subclasses*:

| Java primitive | Writable implementation | Serialized size (bytes) |
|----------------|-------------------------|-------------------------|
| boolean | BooleanWritable | 1 |
| byte | ByteWritable | 1 |
| short | ShortWritable | 2 |
| int | IntWritable | 4 |
| | VIntWritable | 1-5 |
| float | FloatWritable | 4 |
| long | LongWritable | 8 |
| | VLongWritable | 1-9 |
| double | DoubleWritable | 8 |

- > use **get()** and **set()** methods to access the primitive values.



➤ Hadoop IO Data Types (4)



- > **Text** is the implementation for `java.lang.String` class. However some differences:
 - String uses **Unicode chars**, Text uses **UTF-8**. For ASCII differences are unnoticed
 - String `charAt()` returns **char** value, Text `charAt()` returns **int** value
 - String `indexOf()` returns **char** positions, Text `find()` returns **byte** positions
 - String `length()` returns **char** positions, Text `length()` returns **byte** positions
 - For iterations, easier to use String representation than Text
 - Example*:

```
String s = "\u0041\u00DF\u6771\uD801\uDC00";           // s.length():5 , s.indexOf("\u6771"):2 , s.charAt(3): "\uD801"
Text t = new Text("\u0041\u00DF\u6771\uD801\uDC00");   // t.length():10 , t.find("\u6771") :3 , t.charAt(3): 0x6771
```

| Unicode code point | U+0041 | U+00DF | U+6771 | U+10400 |
|---------------------|------------------------|----------------------------|-------------------------------|-------------------------------|
| Name | LATIN CAPITAL LETTER A | LATIN SMALL LETTER SHARP S | N/A (a unified Han Ideograph) | DESERET CAPITAL LETTER LONG I |
| UTF-8 code units | 41 | c3 9f | e6 9d b1 | f0 90 90 80 |
| Java representation | \u0041 | \u00DF | \u6771 | \uD801\uDC00 |

Objectives

- **Mapreduce APIs**
- **Basics**
 - Mapreduce Java API
 - Maven Dependencies
 - WordCount - New Java API
 - WordCount - Old Java API
 - Execution on the cluster
 - MRUnit
 - Hadoop IO data Types (Writable)s
 - **Writing a Writable custom Class**
- **Advanced**
 - How does the data flow through the tasks
- **References and Exercises**



Hadoop IO Data Types (5) - Custom Class Example*



```
import java.io.*;
import org.apache.hadoop.io.*;

public class TextPair implements WritableComparable<TextPair> {
    private Text first; private Text second;

    public TextPair() {
        set(new Text(), new Text());
    }

    public TextPair(String first, String second) {
        set(new Text(first), new Text(second));
    }

    public TextPair(Text first, Text second) {
        set(first, second);
    }

    public void set(Text first, Text second) {
        this.first = first;
        this.second = second;
    }

    public Text getFirst() {
        return first;
    }

    public Text getSecond() {
        return second;
    }
}
```

*Source: Hadoop: The Definitive Guide



Hadoop IO Data Types (6) - Custom Class Example*



```
@Override
public void write(DataOutput out) throws IOException {
    first.write(out);
    second.write(out);
}

@Override
public void readFields(DataInput in) throws IOException {
    first.readFields(in);
    second.readFields(in);
}

@Override
public int hashCode() {
    return first.hashCode() * 163 + second.hashCode();
}

@Override
public boolean equals(Object o) {
    if (o instanceof TextPair) {
        TextPair tp = (TextPair) o;
        return first.equals(tp.first) && second.equals(tp.second);
    }
    return false;
}

@Override
public String toString() {
    return first + "\t" + second;
}
```

*Source: Hadoop: The Definitive Guide



Hadoop IO Data Types (7) - Custom Class Example*



```
@Override
public int compareTo(TextPair tp) {
    int cmp = first.compareTo(tp.first);
    if (cmp != 0) {
        return cmp;
    }
    return second.compareTo(tp.second);
}
```

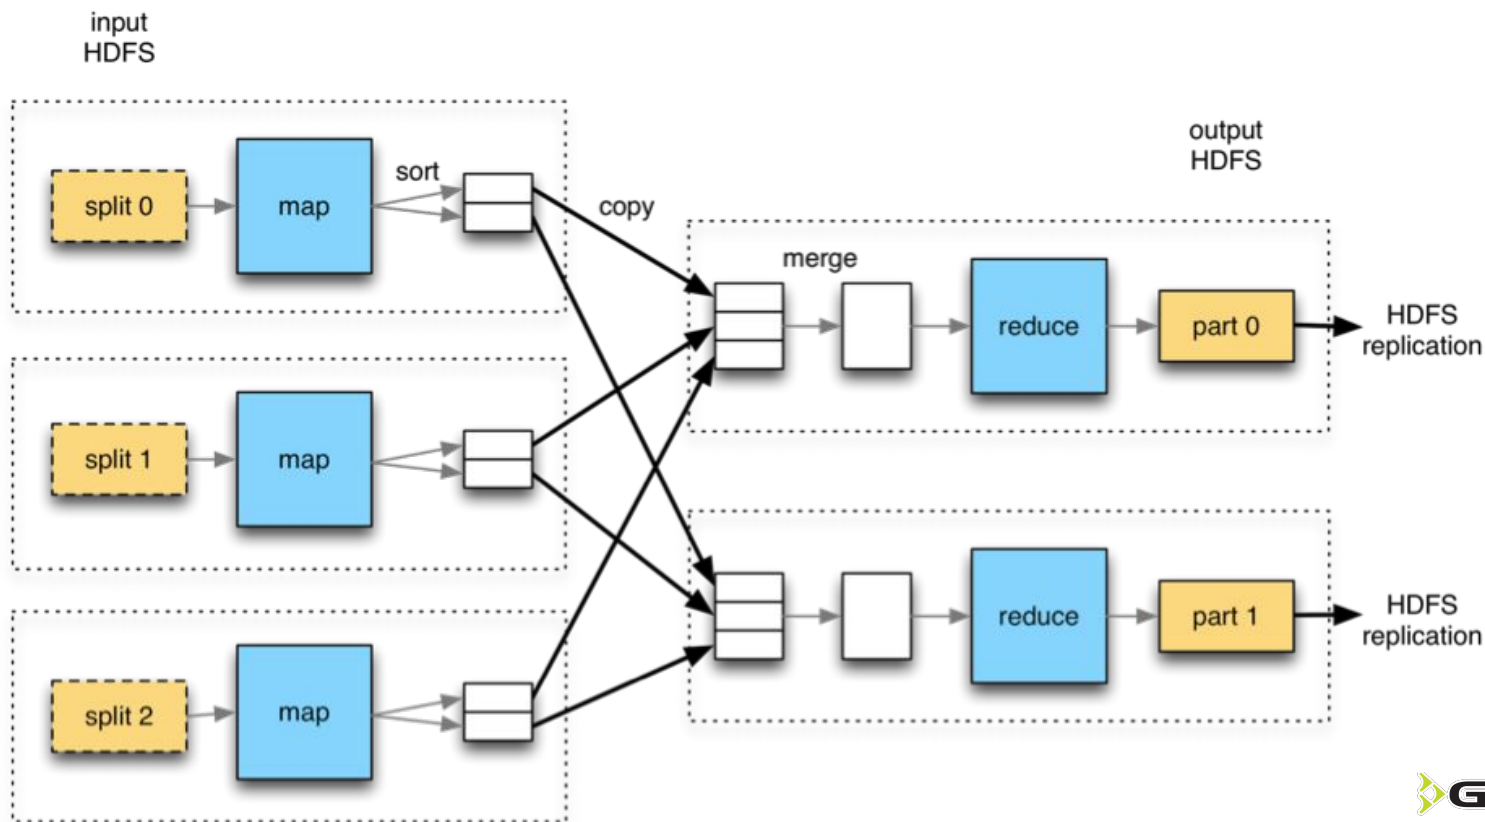
*Source: Hadoop: The Definitive Guide

Objectives

- **Mapreduce APIs**
- **Basics**
 - Mapreduce Java API
 - Maven Dependencies
 - WordCount - New Java API
 - WordCount - Old Java API
 - Execution on the cluster
 - MRUnit
 - Hadoop IO data Types (Writable)s
 - Writing a Writable custom Class
- **Advanced**
 - How does the data flow through the tasks
- **References and Exercises**

MapReduce Java Advanced (1)

How does the data flow through the tasks?



MapReduce Java Advanced (2)



- > Let's considered the default settings of a mapreduce job (in red):

```
public class MinimalMapReduceWithDefaults extends Configured implements Tool {

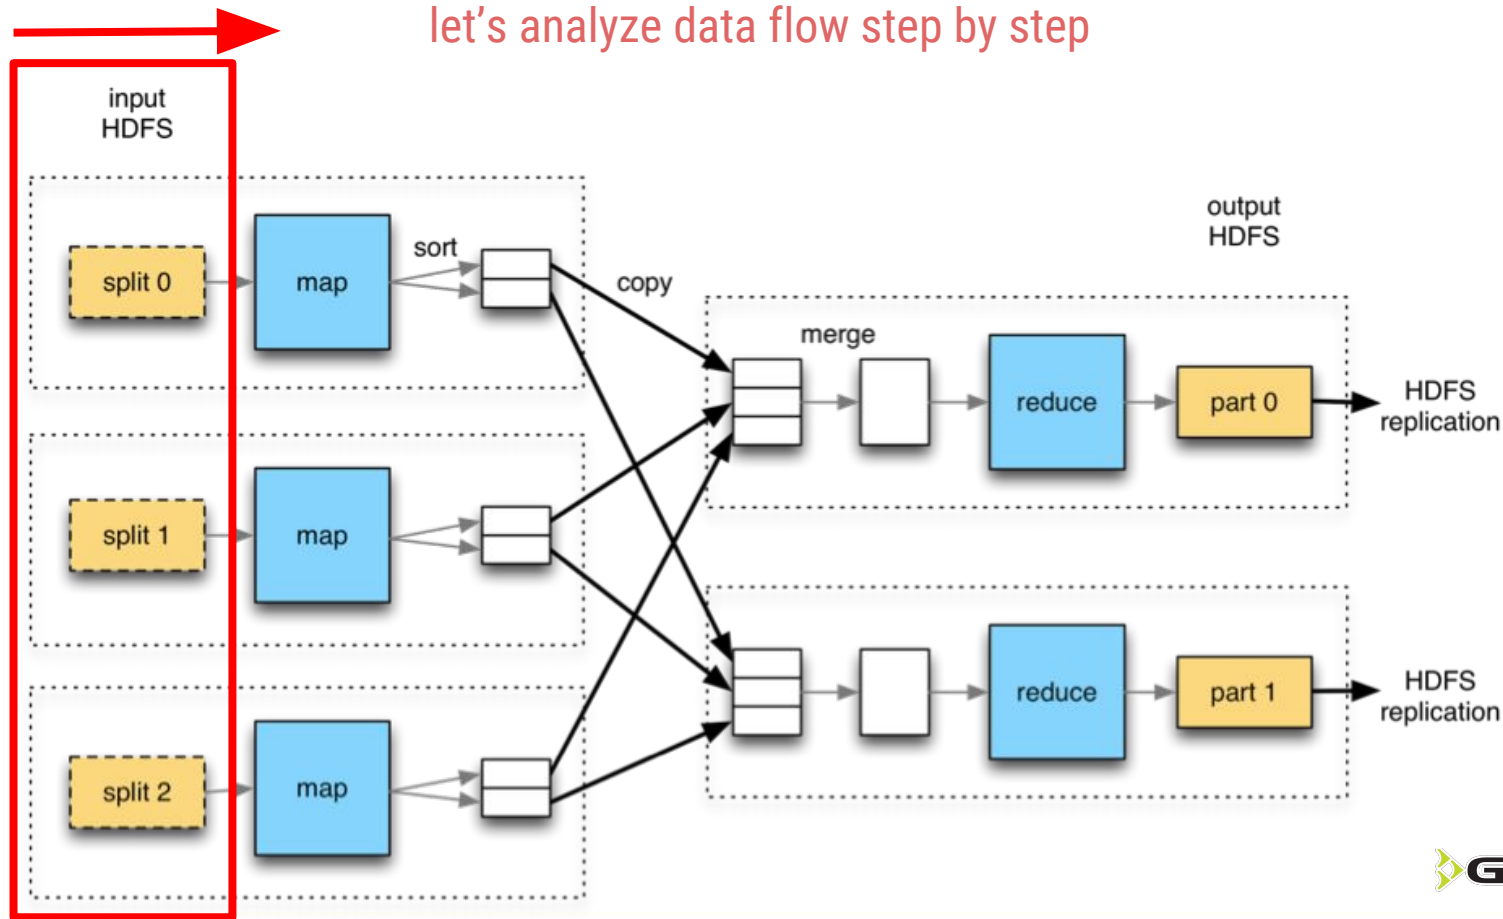
    @Override
    public int run(String[] args) throws Exception {
        Job job = new Job( getConf(), "WordCount" );
        job.setJarByClass( getClass() );
        FileInputFormat.addInputPath( job, new Path(args[0]) );
        FileOutputFormat.setOutputPath( job, new Path(args[1]) );
        job.setInputFormatClass( TextInputFormat.class );
        job.setMapperClass( Mapper.class );
        job.setMapOutputKeyClass( LongWritable.class );
        job.setMapOutputValueClass( Text.class );
        job.setPartitionerClass( HashPartitioner.class );
        job.setNumReduceTasks( 1 );
        job.setReducerClass( Reducer.class );
        job.setOutputKeyClass( LongWritable.class );
        job.setOutputValueClass( Text.class );
        job.setOutputFormatClass( TextOutputFormat.class );
        return job.waitForCompletion(true) ? 0 : 1;
    }

    public static void main(String[] args) throws Exception {
        int exitCode = ToolRunner.run(new MinimalMapReduceWithDefaults(), args);
        System.exit(exitCode);
    }
}
```

*Source: Hadoop: The Definitive Guide

MapReduce Java Advanced (3)

let's analyze data flow step by step



> how the splits work:

```
public class MinimalMapReduceWithDefaults extends Configured implements Tool {

    @Override
    public int run(String[] args) throws Exception {
        Job job = new Job( getConf(), "WordCount" );
        job.setJarByClass( getClass() );
        FileInputFormat.addInputPath( job, new Path(args[0]) );
        FileOutputFormat.setOutputPath( job, new Path(args[1]) );
        job.setInputFormatClass( TextInputFormat.class );
        job.setMapperClass( Mapper.class );
        job.setMapOutputKeyClass( LongWritable.class );
        job.setMapOutputValueClass( Text.class );
        job.setPartitionerClass( HashPartitioner.class );
        job.setNumReduceTasks( 1 );
        job.setReducerClass( Reducer.class );
        job.setOutputKeyClass( LongWritable.class );
        job.setOutputValueClass( Text.class );
        job.setOutputFormatClass( TextOutputFormat.class );
        return job.waitForCompletion(true) ? 0 : 1;
    }

    public static void main(String[] args) throws Exception {
        int exitCode = ToolRunner.run(new MinimalMapReduceWithDefaults(), args);
        System.exit(exitCode);
    }
}
```



- > abstract **InputSplit** represents a Split:

```
public abstract class InputSplit {  
    public abstract long getLength() throws IOException, InterruptedException;  
    public abstract String[] getLocations() throws IOException, InterruptedException;  
}
```

- > Extended class **FileSplit** is used. It represents splits of files with methods like:
 - **getPath()**: local file path.
 - **getStart()**: byte offset of the start of the split from the beginning of the file.
 - **getLength()**: size of the split.



MapReduce Java Advanced (6)

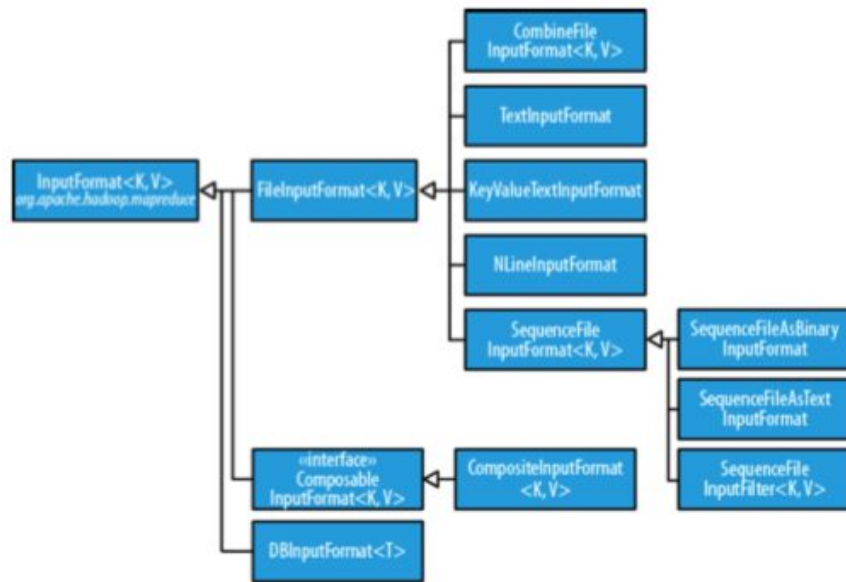


- > abstract **InputFormat** is the base to handle inputs (text files, binary files, databases, etc):

```
public abstract class InputFormat<K, V> {  
    public abstract List<InputSplit> getSplits(JobContext context) throws IOException, InterruptedException;  
    public abstract RecordReader<K, V> createRecordReader(InputSplit split, TaskAttemptContext context) throws IOException,  
    InterruptedException;  
}
```

- > Extended class **FileInputFormat** defines input files in Java Driver
- > 2nd-level extended class **TextInputFormat**:
 - is the **default InputFormat** class for the job.
 - **parses text files** into LongWritable and Text for the mapper.
- > The **client calls getSplits()** of InputFormat to obtain the splits for the job **and sends them to the JobTracker(JT)/ResourceManager(RM)-AppMaster(AM)**
- > The JT/RM-AM **passes the splits** to the TaskTrackers/NodeManager (and they **to the map tasks via context**)

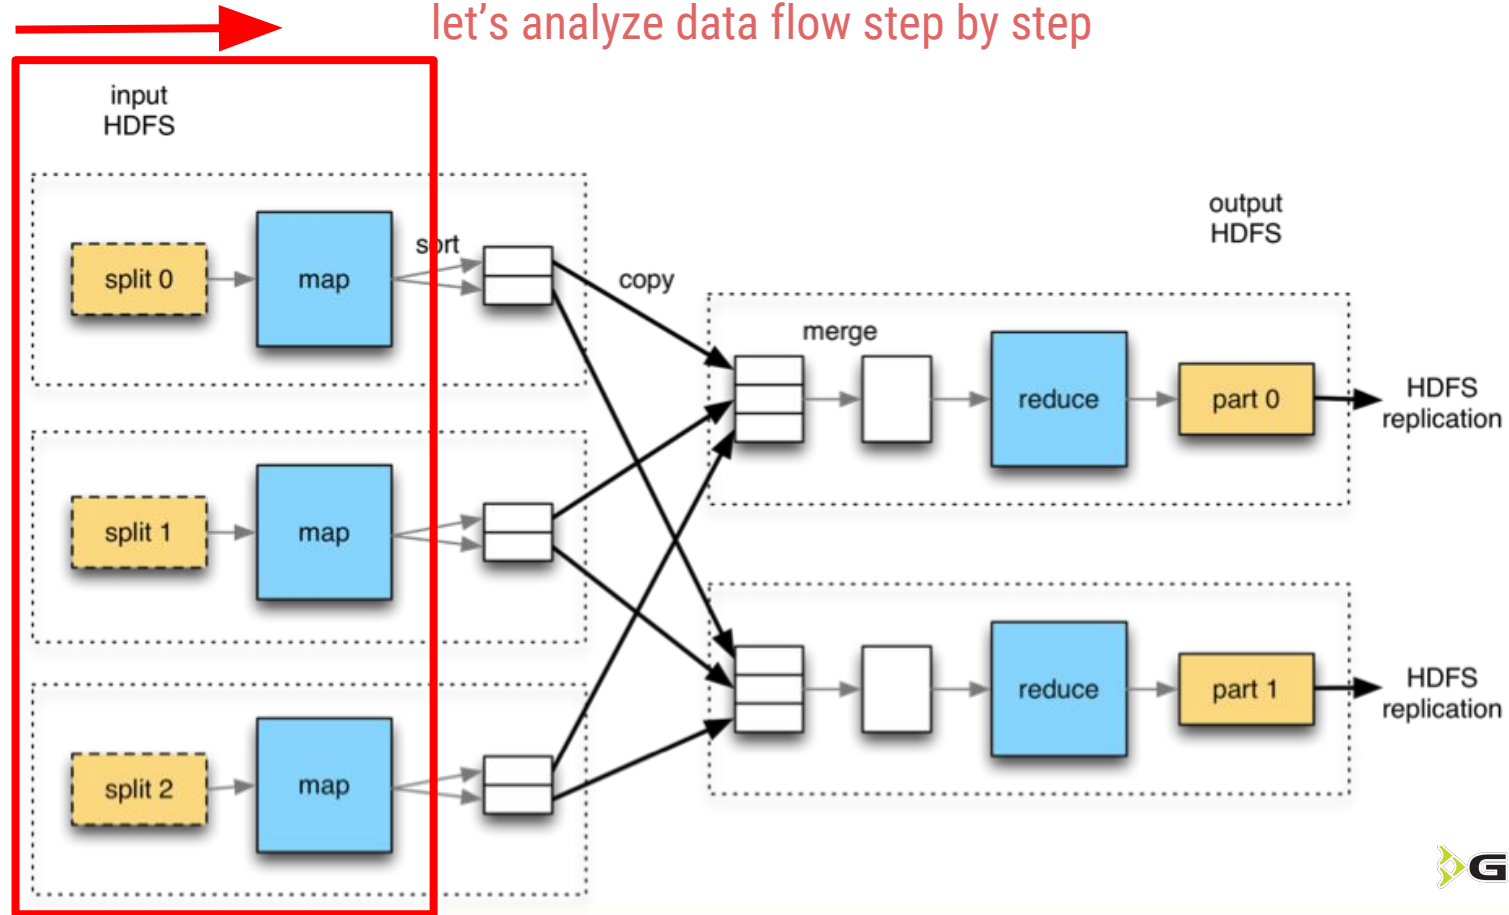
- > Hierarchy of **InputFormat** subclasses*:





MapReduce Java Advanced (8)

let's analyze data flow step by step



> how the map task works:

```
public class MinimalMapReduceWithDefaults extends Configured implements Tool {

    @Override
    public int run(String[] args) throws Exception {
        Job job = new Job( getConf(), "WordCount" );
        job.setJarByClass( getClass() );
        FileInputFormat.addInputPath( job, new Path(args[0]) );
        FileOutputFormat.setOutputPath( job, new Path(args[1]) );
        job.setInputFormatClass( TextInputFormat.class );
        job.setMapperClass( Mapper.class );
        job.setMapOutputKeyClass( LongWritable.class );
        job.setMapOutputValueClass( Text.class );
        job.setPartitionerClass( HashPartitioner.class );
        job.setNumReduceTasks( 1 );
        job.setReducerClass( Reducer.class );
        job.setOutputKeyClass( LongWritable.class );
        job.setOutputValueClass( Text.class );
        job.setOutputFormatClass( TextOutputFormat.class );
        return job.waitForCompletion(true) ? 0 : 1;
    }

    public static void main(String[] args) throws Exception {
        int exitCode = ToolRunner.run(new MinimalMapReduceWithDefaults(), args);
        System.exit(exitCode);
    }
}
```

MapReduce Java Advanced (10)



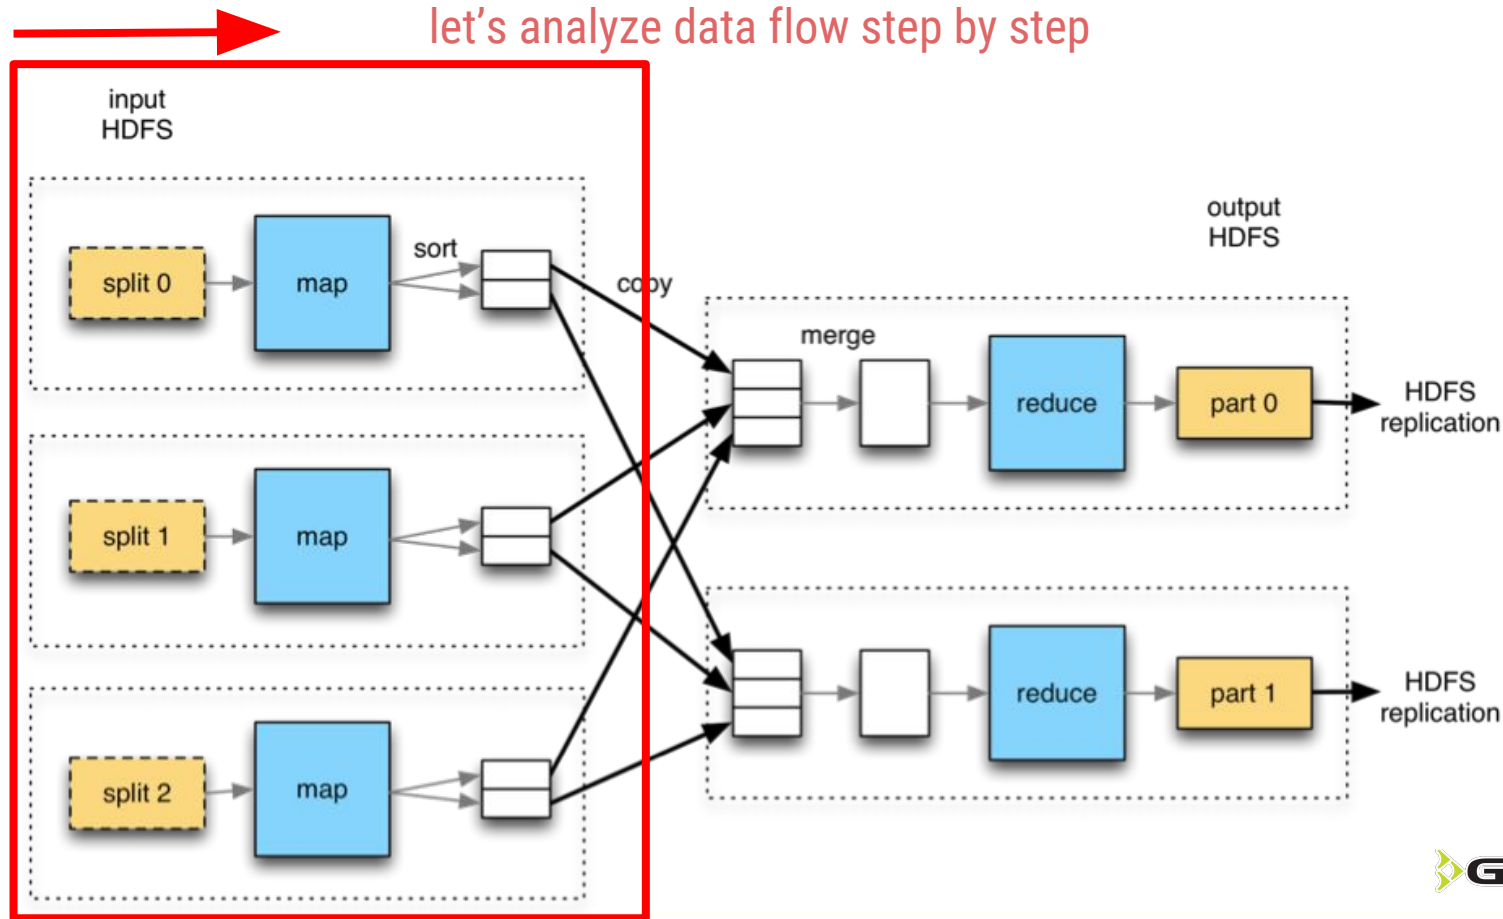
- > abstract **Mapper** implements **run()**:

```
public void run(Context context) throws IOException, InterruptedException {
    setup( context );
    while ( context.nextKeyValue() ) {
        map(context.getCurrentKey(), context.getCurrentValue(), context);
    }
    cleanup(context);
}
```

- > **setup()** method passes the split to **createRecordReader()** of **InputFormat** to open the file and obtain an iterator of records for the part of the file.
- > custom **map()** is called for each of the records passing key, value
- > **cleanup()** closes files, releases resources, etc.

MapReduce Java Advanced (11)

let's analyze data flow step by step



> How the partition works

```
public class MinimalMapReduceWithDefaults extends Configured implements Tool {

    @Override
    public int run(String[] args) throws Exception {
        Job job = new Job( getConf(), "WordCount" );
        job.setJarByClass( getClass() );
        FileInputFormat.addInputPath( job, new Path(args[0]) );
        FileOutputFormat.setOutputPath( job, new Path(args[1]) );
        job.setInputFormatClass( TextInputFormat.class );
        job.setMapperClass( Mapper.class );
        job.setMapOutputKeyClass( LongWritable.class );
        job.setMapOutputValueClass( Text.class );
        job.setPartitionerClass( HashPartitioner.class );
        job.setNumReduceTasks( 1 );
        job.setReducerClass( Reducer.class );
        job.setOutputKeyClass( LongWritable.class );
        job.setOutputValueClass( Text.class );
        job.setOutputFormatClass( TextOutputFormat.class );
        return job.waitForCompletion(true) ? 0 : 1;
    }

    public static void main(String[] args) throws Exception {
        int exitCode = ToolRunner.run(new MinimalMapReduceWithDefaults(), args);
        System.exit(exitCode);
    }
}
```

MapReduce Java Advanced (13)




- > The default partition function:

```
public class HashPartitioner<K, V> extends Partitioner<K, V> {  
    public int getPartition(K key, V value, int numReduceTasks) {  
        return (key.hashCode() & Integer.MAX_VALUE) % numReduceTasks;  
    }  
}
```


MapReduce Java Advanced (14)

let's analyze data flow step by step



> How the reduce task works


```
public class MinimalMapReduceWithDefaults extends Configured implements Tool {

    @Override
    public int run(String[] args) throws Exception {
        Job job = new Job( getConf(), "WordCount" );
        job.setJarByClass( getClass() );
        FileInputFormat.addInputPath( job, new Path(args[0]) );
        FileOutputFormat.setOutputPath( job, new Path(args[1]) );
        job.setInputFormatClass( TextInputFormat.class );
        job.setMapperClass( Mapper.class );
        job.setMapOutputKeyClass( LongWritable.class );
        job.setMapOutputValueClass( Text.class );
        job.setPartitionerClass( HashPartitioner.class );
        job.setNumReduceTasks( 1 );
        job.setReducerClass( Reducer.class );
        job.setOutputKeyClass( LongWritable.class );
        job.setOutputValueClass( Text.class );
        job.setOutputFormatClass( TextOutputFormat.class );
        return job.waitForCompletion(true) ? 0 : 1;
    }

    public static void main(String[] args) throws Exception {
        int exitCode = ToolRunner.run(new MinimalMapReduceWithDefaults(), args);
        System.exit(exitCode);
    }
}
```

MapReduce Java Advanced (16)

let's analyze data flow step by step



MapReduce Java Advanced (15)



>

public class

extends

implements

```
@Override
public int run(String[] args) throws Exception {
    Job job = new Job( getConf(), "WordCount" );
    job.setJarByClass( getClass() );
    FileInputFormat.addInputPath( job, new Path(args[0]) );
    FileOutputFormat.setOutputPath( job, new Path(args[1]) );
    job.setInputFormatClass( TextInputFormat.class );
    job.setMapperClass( Mapper.class );
    job.setMapOutputKeyClass( LongWritable.class );
    job.setMapOutputValueClass( Text.class );
    job.setPartitionerClass( HashPartitioner.class );
    job.setNumReduceTasks( 1 );
    job.setReducerClass( Reducer.class );
    job.setOutputKeyClass( LongWritable.class );
    job.setOutputValueClass( Text.class );
    job.setOutputFormatClass( TextOutputFormat.class );
    return job.waitForCompletion(true) ? 0 : 1;
}

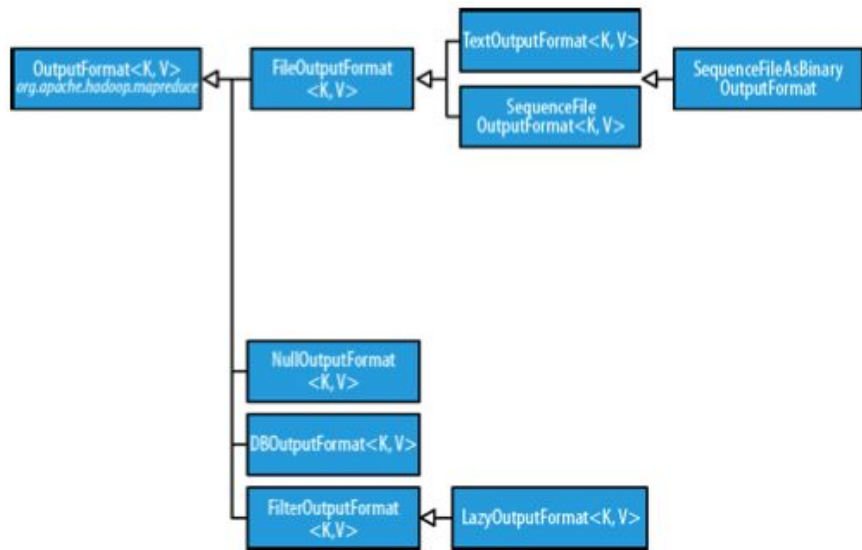
public static void main(String[] args) throws Exception {
    int exitCode = ToolRunner.run(new MinimalMapReduceWithDefaults(), args);
    System.exit(exitCode);
}
}
```

*Source: Hadoop: The Definitive Guide

MapReduce Java Advanced (16)



- > abstract **OutputFormat** has a hierarchy corresponding to **InputFormat**



- > extended **FileOutputFormat** is the default **OutputFormat** class for the job.
- > **FileOutputFormat** writes reducer output key,value pair to part-XXXX file using their **toString()** method

Objectives

- **Mapreduce APIs**
- **Basics**
 - Mapreduce Java API
 - Maven Dependencies
 - WordCount - New Java API
 - WordCount - Old Java API
 - Execution on the cluster
 - MRUnit
 - Hadoop IO data Types (Writable)s
 - Writing a Writable custom Class
- **Advanced**
 - How does the data flow through the tasks
- **References and Exercises**



Mapreduce - References



- > Includes some documentation about the old api

http://hadoop.apache.org/docs/r1.2.1/mapred_tutorial.html

- > *White, Tom*. Hadoop: The Definitive Guide



Mapreduce - Exercise 1



Inverted Index

Given a directory with books in txt format, write a mapreduce which outputs an inverted index, i.e., a table that associates a word with the books and the corresponding positions at which it occurs (http://en.wikipedia.org/wiki/Inverted_index).

Dataset URL: [here](#)

HDFS Cluster DataSet path: `/user/hadoop/mapreduce/data/books`

hint 1: Suggested output example (not real data):

Love `alice_in_wonderland.txt:100,the_prince.txt:900,the_prince.txt:1050`

hint 2: The mapper doesn't receive explicitly the filename as input, but the context has access to the `InputSplit` and casting it to a `FileSplit` object, the `getPath()` method can be used.



Mapreduce - Exercise 2



Column-wise Variance (s^2) of a matrix

Given a csv file without headers, calculate the sample variance (s^2) of each column.

(<http://en.wikipedia.org/wiki/Variance>)

HDFS DataSet path: `/user/hadoop/mapreduce/data/matrix`

!!!DON'T DOWNLOAD IT FROM THE CLUSTER \$\$\$\$!!!

$$s^2 = \frac{1}{(N-1)} \sum_{i=1}^N (x_i - \bar{x})^2$$

hint 1: Suggested output: `columnIndex<tab>sampleVariance`.

Example:

```
0    135.6
1     2.2
2   536.9
...
```

hint 2: Assume the file has only numeric values and no entries are missing (no NULLs or empty).

hint 3: Focus on the mapreduce. (If any) other plumbing code is required, do it manually, afterwards implement it if you have the time (preferably with python).



Mapreduce - Exercise 2



Column-wise Variance (s^2) of a matrix

question 1:

Assuming your matrix is very big (both rows and columns), how many reducers does it make sense to have?

question 2:

How many mapreduce jobs do you need?

$$s^2 = \frac{1}{(N-1)} \sum_{i=1}^N (x_i - \bar{x})^2$$

