



# Are you ready for the future?

High Performance Solutions, Backend Development and Integration Services, Embedded  
Software Development, Big Data (Visualization, Architecture, Science), Business  
Intelligence & Analytics



**The following presentation is based on  
the Apache Hive Documentation,  
Programming Hive book, Hadoop the Definitive Guide  
and the experience of some developers from  
the Big Data Studio  
for internal training purposes.**





# Hive - What it is

## What it is

- Structured Data Warehouse over Hadoop
- SQL -> Map Reduce
- HDFS Storage
- Metadata in RDBMS
- Aggregation functions

## What it is NOT

- SchemalessDB
- TransactionalDB
- Real Time platform



# Hive - What it is

## What it is NOT



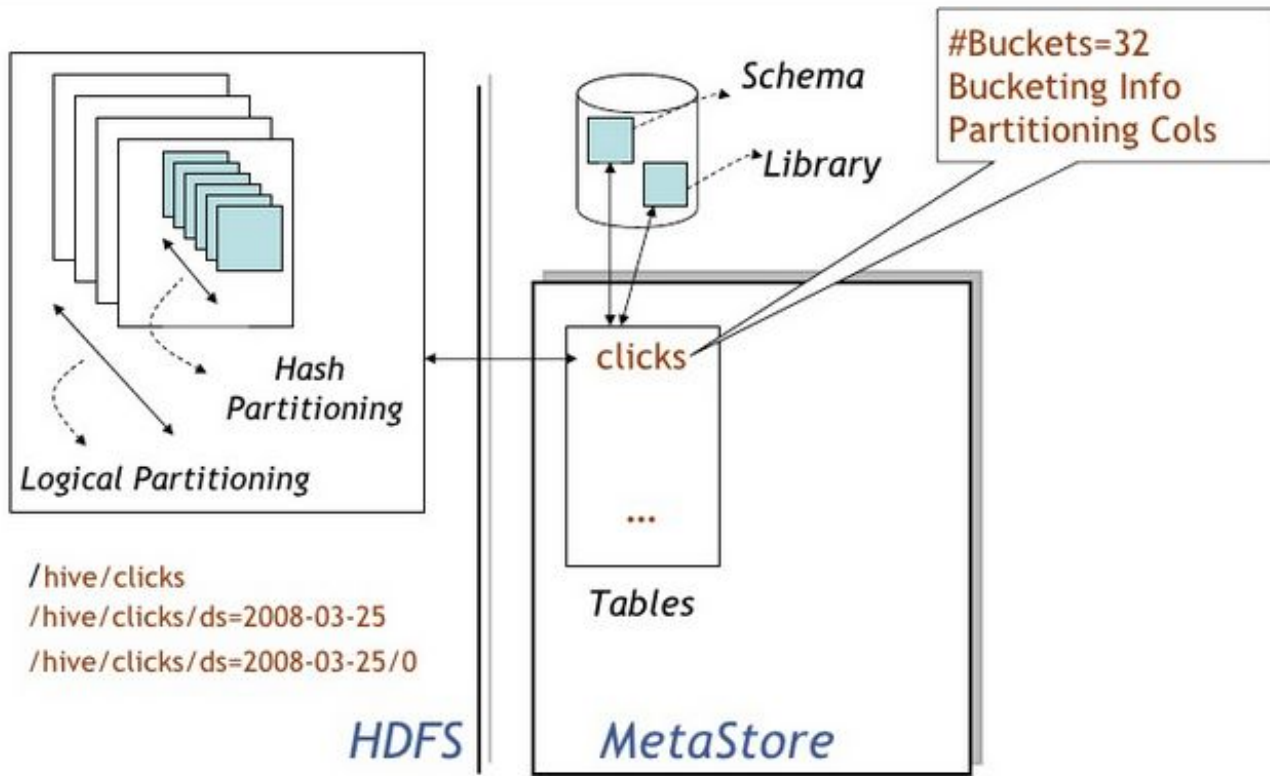
Hive is **not** designed for online transaction processing (OLTP) workloads.

It is best used for traditional data warehousing tasks.



# Hive - Architecture

## Data Model





# Hive - Shell Client (hive cli)

---

## 1) Interactive:

```
hive
```

## 2) Non Interactive:

```
hive -S -e "select * from ventas ;"
```

```
-S -v -f -e --hivevar --hiveconf
```



# Hive - Data Types: Basic

---

**TINYINT:** 1 byte signed integer.

**SMALLINT:** 2 byte signed integer.

**INT:** 4 byte signed integer.

**BIGINT:** 8 byte signed integer.

**BOOLEAN:** Boolean true or false.

**FLOAT:** Single precision floating point.

**DOUBLE:** Double precision floating point.

**STRING:** Sequence of characters.

**TIMESTAMP:** Integer, float, or string. Compatible con java.sql.Timestamp

**BINARY:** Array of bytes.





# Hive - Data Types: Complex

---

## ➤ ARRAY array('Juan', 'Luis')

Ordered sequences of the same type that are indexable using zero-based integers. For example, if a column name is of type ARRAY of strings with the value ['John', 'Doe'], then the second element can be referenced using name[1].

## ➤ MAP map('Nombre': 'Juan', 'Apellido': 'Perez')

A collection of key-value tuples, where the fields are accessed using array notation (e.g., ['key']). For example, if a column name is of type MAP with key→value pairs 'first'→'John' and 'last'→'Doe', then the last name can be referenced using name['last'].

## ➤ STRUCT: ('Nombre':'Juan', 'Edad':40, 'Sexo':'M')

Analogous to a C struct or an “object.” Fields can be accessed using the “dot” notation. For example, if a column name is of:

**type STRUCT:name {first STRING; last STRING}, then**

the first name field can be referenced using name.first.



# Hive - Schema On Read

## Schema on WRITE



1 Create schema

2 Load Data

3 Read Data



Stuck with one-size-fits all schema



Complex with multiple-datasets



Bad for Unknown datasets

## Schema on READ



1 Load Data as-is

2 Create Schema

3 Read it



If the schema doesn't match, it fills the field with null values



Useful to work with multiple datasets




Good for know datasets



# Internal Tables

---



**Hive controls the lifecycle of their data.**



**Hive stores the data for these tables in a subdirectory under the directory defined by `hive.metastore.warehouse.dir`**




**Less convenient for shared data between different tools**



When we **DROP** an internal table  
**Hive deletes the data  
inside the table**



# External Tables



- Hive does not assume it owns the data.
- **Dropping the table does not delete the data,**
- (although the metadata for the table will be deleted)

```
CREATE EXTERNAL TABLE IF NOT EXISTS stocks (  
  exchange      STRING,  
  symbol        STRING,  
  ymd           STRING,  
  price_open    FLOAT,  
  price_high    FLOAT,  
  price_low     FLOAT,  
  price_close   FLOAT,  
  volume        INT,  
  price_adj_close FLOAT)  
ROW FORMAT DELIMITED FIELDS TERMINATED BY ','  
LOCATION '/data/stocks';
```



# Partitioned Tables

```
CREATE EXTERNAL TABLE IF NOT EXISTS log_messages (  
  hms          INT,  
  severity     STRING,  
  server       STRING,  
  process_id   INT,  
  message      STRING)  
PARTITIONED BY (year INT, month INT, day INT)  
ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t';
```

```
hdfs://master/user/hive/log_messages  
./employees/year=2012/month=October/day=12  
./employees/year=2012/month=November/day=12
```

Users of the table don't need to care if these "columns" are partitions or not, except when they want to **optimize query performance**

If we add new data to a folder declared as partition, that data will appear available for the table after the insertion

You can declare a folder as a partition.

A query without a WHERE clause will query all the data in all the partitions.  
If you add the WHERE clause, it will read only the partition (which is translated as accessing a folder) that match with that condition.

However, the partition you create makes a pseudocolumn on which you can query



# Bucketed Tables

```
CREATE TABLE bucketed_users (  
  id INT,  
  name STRING)  
CLUSTERED BY (id) INTO 4 BUCKETS;  
  
set hive.enforce.bucketing= true;
```

```
hdfs://master/user/hive/bucketed_users  
./00000_0  
./00000_1  
./00000_2  
./00000_3
```

Data is stored in as many files as buckets inside the partition/table folder.

User unaware if table is bucketed or not

+ Bucketed tables are fantastic in that they allow much more efficient **sampling**

+ Allows **efficient queries**, e.g., “map-side join”

A Clustered Field value is assigned a bucket, the Record is put on the file of that bucket (similar to HashPartitioner.class in mapreduce)

**set** hive.enforce.bucketing= **true**;  
sets same number of reducers as buckets during inserts to table.



# Bucketed Tables

---

```
CREATE TABLE page_view(viewTime INT, userid
BIGINT,
    page_url STRING, referrer_url STRING,
    ip STRING COMMENT 'IP Address of the
User')
COMMENT 'This is the page view table'
PARTITIONED BY(dt STRING, country STRING)
CLUSTERED BY(userid) SORTED BY(viewTime)
INTO 32 BUCKETS
ROW FORMAT DELIMITED
    FIELDS TERMINATED BY '\001'
    COLLECTION ITEMS TERMINATED BY '\002'
    MAP KEYS TERMINATED BY '\003'
STORED AS SEQUENCEFILE;
```

The page\_view table is bucketed (clustered by) *userid* and within each bucket the data is sorted in increasing order of viewTime. Such an organization allows the user to **do efficient sampling** on the clustered column - in this case userid

.....

The CLUSTERED BY and SORTED BY creation commands do not affect how data is inserted into a table – only **how it is read**.

.....

Users must be careful to insert data correctly by specifying the number of reducers to be equal to the number of buckets, and using CLUSTER BY and SORT BY commands in their query.





# Temporary Tables & Views

## TEMPORARY TABLES

- Will only be visible to the current session
- No support for creation of indexes
- Partition columns are not supported.

## VIEWS

- A purely logical object with no associated storage
- A view's schema is frozen at the time the view is created; subsequent changes to underlying tables (e.g. adding a column) will not be reflected in the view's schema.
- Views are read-only



## Hive - Reading files...Row way

---

```
CREATE TABLE usuarios(  
    nombre STRING,  
    edad INT)  
ROW FORMAT DELIMITED  
    FIELDS TERMINATED BY '\t'  
    LINES TERMINATED BY '\n';
```



# Hive - Reading file...SerDe way

- HDFS Files → InputFileFormat →  
<key, value> → Deserializer  
→ Row object
  - Example: Query data
- Row object → Serializer → <key, value>  
→ OutputFileFormat → HDFS Files
  - Insert

```
CREATE EXTERNAL TABLE tweets (  
  ...  
  retweeted_status STRUCT<  
    text:STRING,  
    user:STRUCT<  
      screen_name:STRING,  
      name:STRING  
    >  
  >,  
  text STRING,  
  ...  
)  
ROW FORMAT SERDE  
  'com.cloudera.hive.serde.JSONSerDe'
```



# Hive - Reading files...Serde way

---

Using Json SerDe you can save storage

*If your schema has a lot of columns and often they are null. It's a good idea to use SerDe. You can create json without the keys (columns) and Serde will show it as null.*



# Loading Data into a table

```
LOAD DATA [LOCAL] INPATH 'filepath' [OVERWRITE]
INTO TABLE tablename [PARTITION (partcol1=val1,
partcol2=val2 ...)]
```

In this case, load operations are currently **pure copy/move operations** that move datafiles into locations corresponding to Hive tables.

Since Hive 0.14:

```
INSERT INTO TABLE tablename [PARTITION
(partcol1[=val1], partcol2[=val2] ...)] VALUES
values_row [, values_row ...]
```

**Compressed Data Storage** : Keeping data compressed in Hive tables has, in some cases, been known to give **better performance** than uncompressed storage; both in terms of **disk usage** and **query performance**.

.....  
You can import text files compressed with Gzip or Bzip2 directly into a table stored as TextFile

.....  
**However, in this case Hadoop will not be able to split your file into chunks/blocks and run multiple maps in parallel.**

# Loading Data into a table with another format

---

```
CREATE TABLE raw (line STRING)
  ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t' LINES TERMINATED BY '\n';

CREATE TABLE raw_sequence (line STRING)
  STORED AS SEQUENCEFILE;

LOAD DATA LOCAL INPATH '/tmp/weblogs/20090603-access.log.gz' INTO TABLE raw;

SET hive.exec.compress.output=true;
SET io.seqfile.compression.type=BLOCK; -- NONE/RECORD/BLOCK (see below)
INSERT OVERWRITE TABLE raw_sequence SELECT * FROM raw;
```



# Inserting Data into a table

```
//Overwrite data in tablename1
INSERT OVERWRITE TABLE tablename1 [PARTITION
(partcol1=val1, partcol2=val2 ...)] [IF NOT
EXISTS]]
select_statement1 FROM from_statement;

// Append the data to tablename1
INSERT INTO TABLE tablename1
[PARTITION (partcol1=val1, partcol2=val2
...)]
select_statement1 FROM from_statement;
```

**Static Partition Insert** : Define partitions in the insert query.

```
FROM staged_employees se
INSERT OVERWRITE TABLE employees
  PARTITION (country = 'US', state = 'OR')
  SELECT * WHERE se.cnty = 'US' AND se.st = 'OR'
INSERT OVERWRITE TABLE employees
  PARTITION (country = 'US', state = 'CA')
  SELECT * WHERE se.cnty = 'US' AND se.st = 'CA'
INSERT OVERWRITE TABLE employees
  PARTITION (country = 'US', state = 'IL')
  SELECT * WHERE se.cnty = 'US' AND se.st = 'IL';
```

**Dynamic Partition Insert** : Infer the partitions to create based on query parameters.

```
INSERT OVERWRITE TABLE employees
PARTITION (country, state)
SELECT ..., se.cnty, se.st
FROM staged_employees se;
```



# Inserting Data into a table - Create Table as Select (CTAS)

The table created by CTAS is **atomic**, meaning that the table is not seen by other users until all the query results are populated.

CTAS has these restrictions:

- The target table cannot be a partitioned table.
- The target table cannot be an external table.
- The target table cannot be a list bucketing table.


## Example:

```
CREATE TABLE new_key_value_store
  ROW FORMAT SERDE "org.apache.hadoop.hive.serde2.columnar.ColumnarSerDe"
  STORED AS RCFile
  AS
SELECT (key % 1024) new_key, concat(key, value) key_value_pair
FROM key_value_store
SORT BY new_key, key_value_pair;
```





# HiveQL: Queries - SELECT





# HiveQL: Queries - SELECT

```
> SELECT [ALL | DISTINCT] select_expr, select_expr, ...  
FROM table_reference  
[WHERE where_condition]  
[GROUP BY col_list]  
[CLUSTER BY col_list | [DISTRIBUTE BY col_list] [SORT BY col_list]]  
[LIMIT number];  
  
> SELECT * FROM log_messages WHERE year = 2014 AND month = 'December';  
  
> SELECT name FROM employees WHERE surname LIKE '%Mess.';
```



# HiveQL: Queries - JOINS I

USE **EXPLAIN**  
TO CHECK  
PERFORMANCE

```
> SELECT a.* FROM a JOIN b ON (a.id = b.id AND a.department = b.department);
```

```
> NOT ALLOWED(NON-Equality): SELECT a.* FROM a JOIN b ON (a.id <> b.id);
```

#Hive converts joins over multiple tables into a single map/reduce job if for every table the same column is used in the join clause. if not, it is converted into multiple map/reduce jobs

```
> SELECT a.val, b.val, c.val FROM a JOIN b ON (a.key = b.key1) JOIN c ON (c.key = b.key1) #1  
M-R
```

#OUTER JOINS: Outer join allow you to find nonmatches in the tables being joined.

```
> SELECT sales.* FROM sales RIGHT OUTER JOIN employee ON sales.employeeid = employee.id
```

## #SEMI JOINS

```
> SELECT * FROM singers WHERE singers.id IN (SELECT sales.singerId from sales)
```



# HiveQL: Queries - JOINS II

USE **EXPLAIN**  
TO CHECK  
PERFORMANCE

**#MAP JOINS:** If one table fits in memory, Hive can load it into memory in the map phase. It does not need a reducer.

```
> SELECT /*+ MAPJOIN(b) */ a.key, a.value FROM a JOIN b ON a.key = b.key;
```

```
> SELECT /*+ MAPJOIN(c) */ * FROM orders o JOIN cities c ON (o.city_id = c.id);
```

(!) Bucketing:

Map joins can take advantage of bucketed tables, since a mapper working on a bucket of the left table needs to load only the corresponding buckets of the right table to perform the join. The syntax for the join is the same as for the in-memory case shown earlier; however, you also need to enable the optimization with the following:

```
SET hive.optimize.bucketmapjoin=true;
```



# HiveQL: Queries - Subqueries

#Hive has limited support for subqueries, permitting a subquery in the **FROM** clause, in a **WHERE** clause or inside an **IN** or **EXISTS** statement.

```
SELECT t3.col  
FROM (  
  SELECT a+b AS col  
  FROM t1  
  UNION ALL  
  SELECT c+d AS col  
  FROM t2  
) t3
```

```
SELECT A  
FROM T1  
WHERE EXISTS (  
  SELECT B FROM T2  
  WHERE T1.X = T2.Y)
```

```
SELECT *  
FROM A  
WHERE A.a IN  
(SELECT foo FROM B);
```



# HiveQL: ORDER BY

# **ORDER BY**: it performs a total Ordering of the query result set. This means that all the data is passed through a *single* reducer. **May take long time to execute.**

```
> SELECT s.year, s.id, s.price FROM stock s ORDER BY s.year ASC, s.price DESC
```

# **SORT BY**: it performs a local Ordering of the query result set. This means that all the data is passed through a *single* reducer. **May take long time to execute.**

```
> SELECT s.year, s.id, s.price FROM stock s SORT BY s.year ASC, s.price DESC
```

# **DISTRIBUTED BY WITH SORT BY**: controls how map output is divided among reducers.

# We can use **DISTRIBUTE BY** to ensure that the records for each stock symbol go to the same reducer, then use **SORT BY** to order the data the way we want.

```
> SELECT s.ymd, s.symbol, s.price_close
```

```
> FROM stocks s
```

```
> DISTRIBUTE BY s.symbol
```

```
> SORT BY s.symbol ASC, s.ymd ASC;
```



# Hive: INDEXES

```
CREATE INDEX index_name
ON TABLE base_table_name (col_name, ...)
AS 'index.handler.class.name'
[WITH DEFERRED REBUILD]
[IDXPARTITION (property_name=property_value, ...)]
[IN TABLE index_table_name]
[PARTITIONED BY (col_name, ...)]
[
  [ ROW FORMAT ...] STORED AS ...
  | STORED BY ...
]
[LOCATION hdfs_path]
[TBLPROPERTIES (...)]
[COMMENT "index comment"]
```

- ❑ Improve the speed of query lookup on certain columns of a table
- ❑ Additional processing to create the index and disk space to store the index
- ❑ Only single-table indexes are supported.
- ❑ By default, index partitioning matches the partitioning of the base table. For example, a table may be partitioned by date+region even though the index is partitioned by date alone



# Hive: Compression & FileFormat

## Compression

- ↑ Save disk storage. Reduce traffic in I/O.
- ↓ CPU Overhead.

Trade-Off between Compression Codes:

-> Speed to de/compress -> Space on Disk

(!) Important: Splittable

## Compression Codecs

- LZ0
- Snappy
- Bzip2 (!)
- LZ0(!)

## File Formats:

- ❑ Default: Text format
- ❑ Sequence Files
- ❑ RCFile

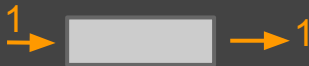
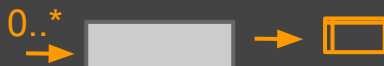
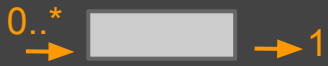
### CONVERT FROM ONE FORMAT TO ANOTHER:

When a query SELECTs from one table and INSERTs into another, Hive uses the metadata about the tables and handles the conversion automatically.





# Hive - Functions

- Standard Functions (UDF, or just DF) 
- Table Generating Function (UDTF) 
- Aggregation Function (UDAF) 

## ➤ Macros

```
hive> CREATE TEMPORARY MACRO SIGMOID (x DOUBLE) 1.0 / (1.0 + EXP(-x));  
hive> SELECT SIGMOID(2) FROM src LIMIT 1;
```



# Hive - Meta comandos

**> show functions;**

**List functions and available operations**

**> describe function [function name]**

**Shows a description of the function**

**> describe function extended [function name]**

**Show a detailed description of the function**



**Presto is a distributed SQL query engine optimized for ad-hoc analysis at interactive speed. It supports standard ANSI SQL, including complex queries, aggregations, joins, and window functions.**

**The execution model of Presto is fundamentally different from Hive/MapReduce. Hive translates queries into multiple stages of MapReduce tasks that execute one after another. Each task reads inputs from disk and writes intermediate output back to disk. In contrast, the Presto engine does not use MapReduce. It employs a custom query and execution engine with operators designed to support SQL semantics. In addition to improved scheduling, all processing is in memory and pipelined across the network between stages. This avoids unnecessary I/O and associated latency overhead. The pipelined execution model runs multiple stages at once, and streams data from one stage to the next as it becomes available. This significantly reduces end-to-end latency for many types of queries.**



# Hive - Good Practices

---

Partitions	Large files rather than smaller ones
De-normalization	Strict Mode
Compression	Use indexes
Parallelization	Reuse JVM
Optimize Joins	Be carefull with the sintax in your queries.
Tune number of reducers	

*Source- Hive Programming chapter 10*



# Hive - Exercise

## Column-wise Variance ( $s^2$ ) of a matrix

Given a csv file without headers, calculate the sample variance ( $s^2$ ) of each column.

(<http://en.wikipedia.org/wiki/Variance>)

HDFS DataSet path: `/user/hadoop/mapreduce/data/matrix`

!!!DON'T DOWNLOAD IT FROM THE CLUSTER \$\$\$\$!!!

$$s^2 = \frac{1}{(N-1)} \sum_{i=1}^N (x_i - \bar{x})^2$$

**hint 1:** Suggested output: `sampleVariance1<tab>sampleVariance2<tab>sampleVariance3...`

Example:

`135.6 2.2 536.9 ...`

**hint 2:** Assume the file has only numeric values and no entries are missing (no NULLs or empty).

# References

We create innovative software products that  
appeal to global audiences.





## References

---

- **Apache Hive Documentation.**- <https://cwiki.apache.org/confluence/display/Hive/Home>
- **Hadoop: The Definitive Guide, 2nd Edition** (Chapter 11). O'Reilly Media / Yahoo Press - [Online @ Globant's Big Data Training Site](#)
- **Programming Hive.** O'Reilly Media - [Online @ Globant's Big Data Training Site](#)

**Thanks**